

Request-Oriented Durable Write Caching for Application Performance

Sangwook Kim¹, Hwanju Kim², Sang-Hoon Kim³, Joonwon Lee¹,
and Jinkyu Jeong¹


Sungkyunkwan University¹

University of Cambridge²

Korea Advanced Institute of Science and Technology (KAIST)³

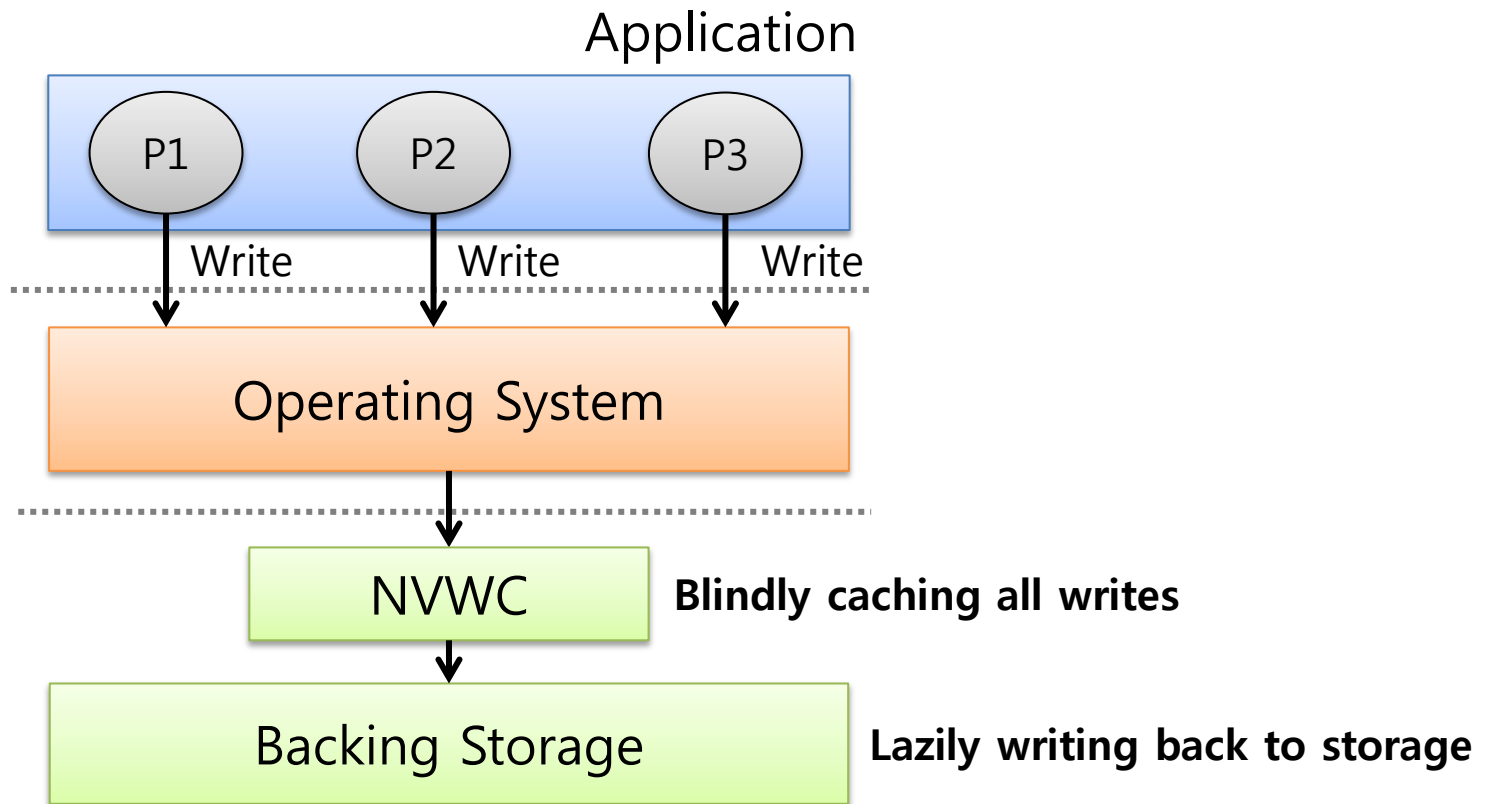
Non-volatile Write Cache

- **Volatile** DRAM cache is ineffective for write
 - Writes are dominant I/Os [FAST'09, FAST'10, FAST'14]
- Non-volatile write cache (NVWC) provides
 - **Fast response** for write w/o loss of durability
- NVWC candidates



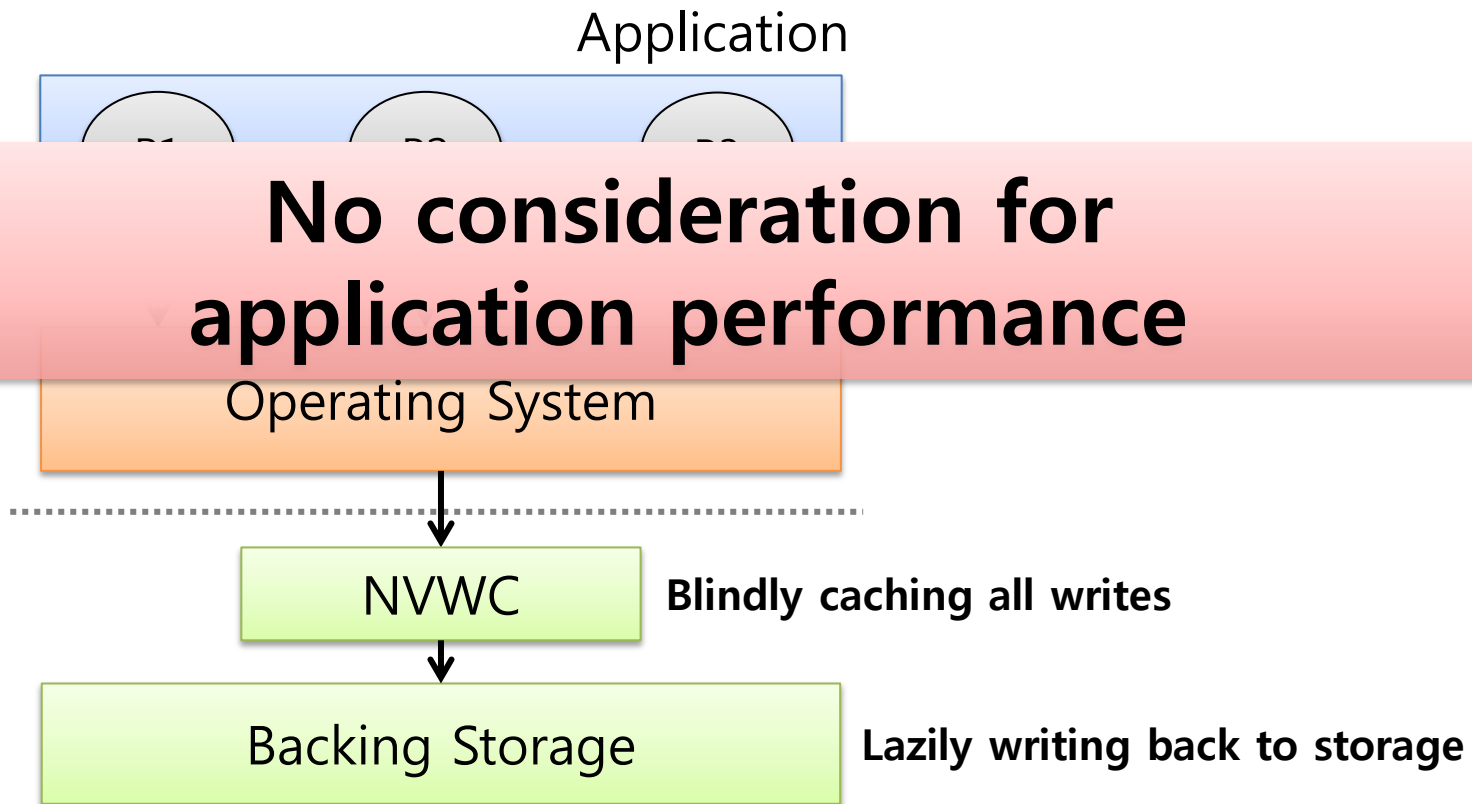
Non-volatile Write Cache

- Simple caching policy



Non-volatile Write Cache

- Simple caching policy

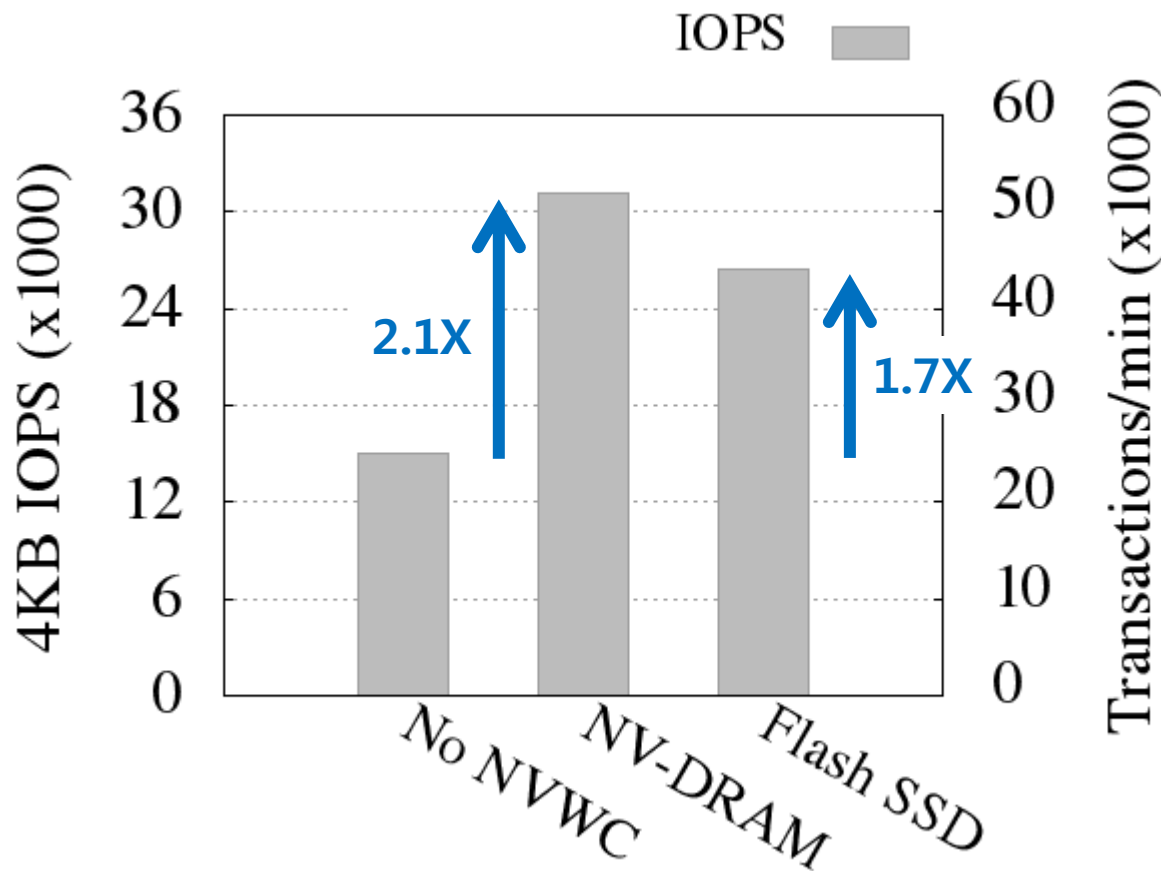


Impact on Application Performance

- Illustrative experiment
 - TPC-C workload
 - PostgreSQL database
 - 2 NVWC devices
 - 32MB NV-DRAM (emulated via ramdisk)
 - 4GB flash SSD

Impact on Application Performance

- Experimental result

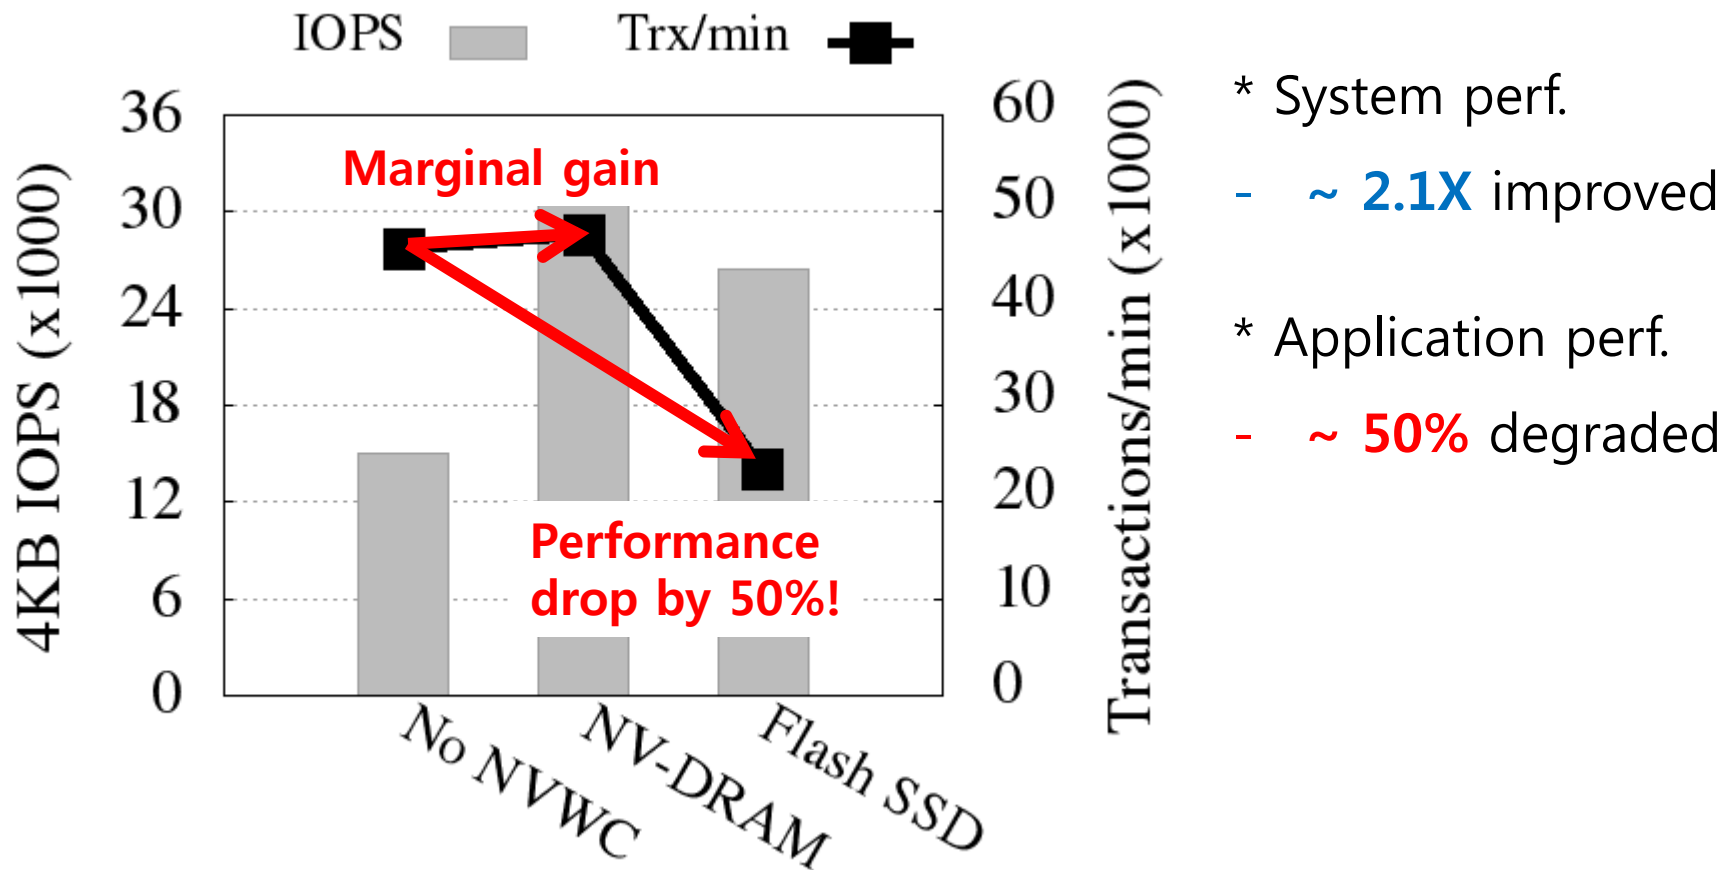


* System perf.

- ~ **2.1X** improved

Impact on Application Performance

- Experimental result



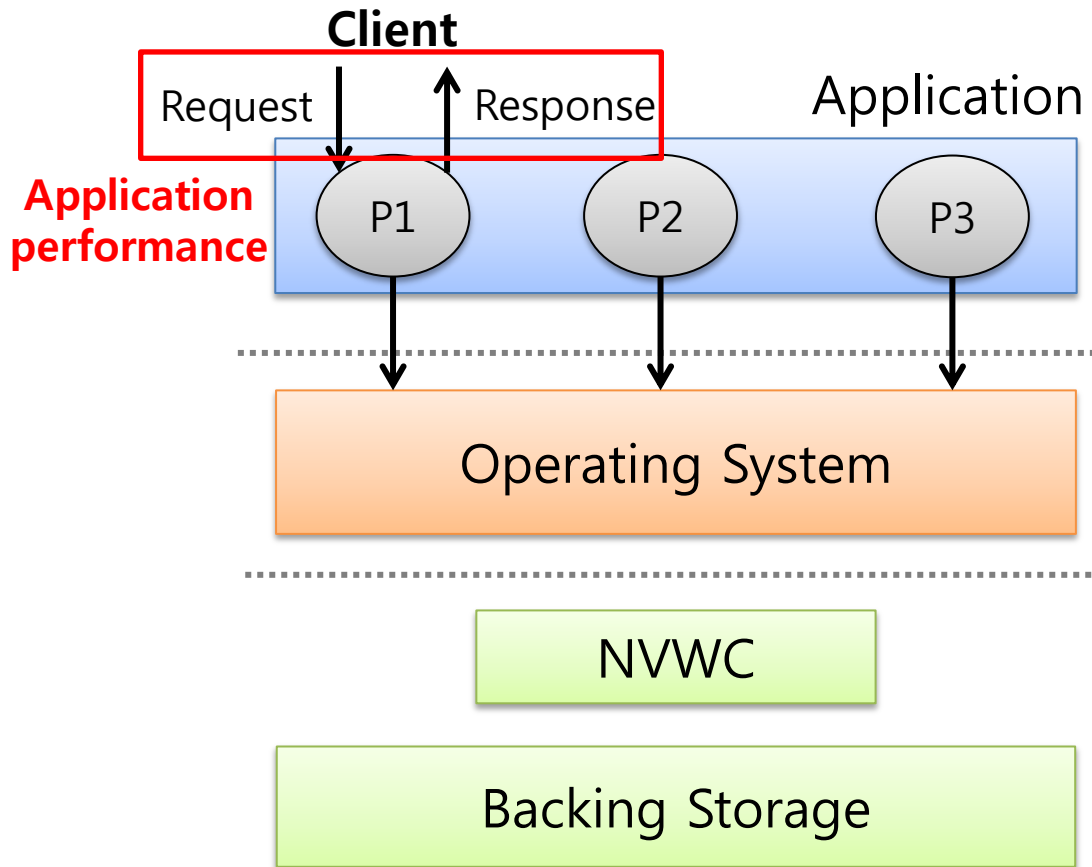
What's the Problem?

- **Criticality-agnostic** contention



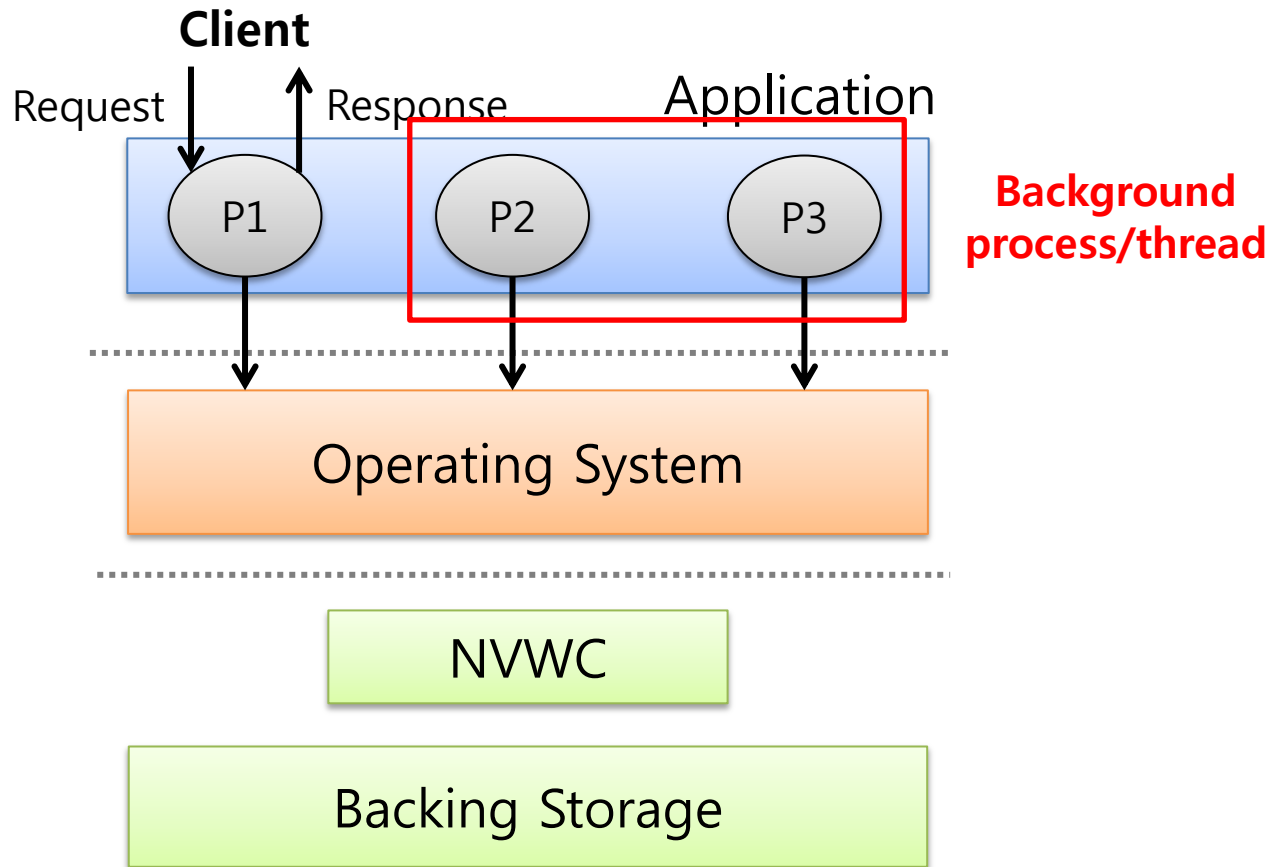
Criticality-Agnostic Contention

- Different write criticality



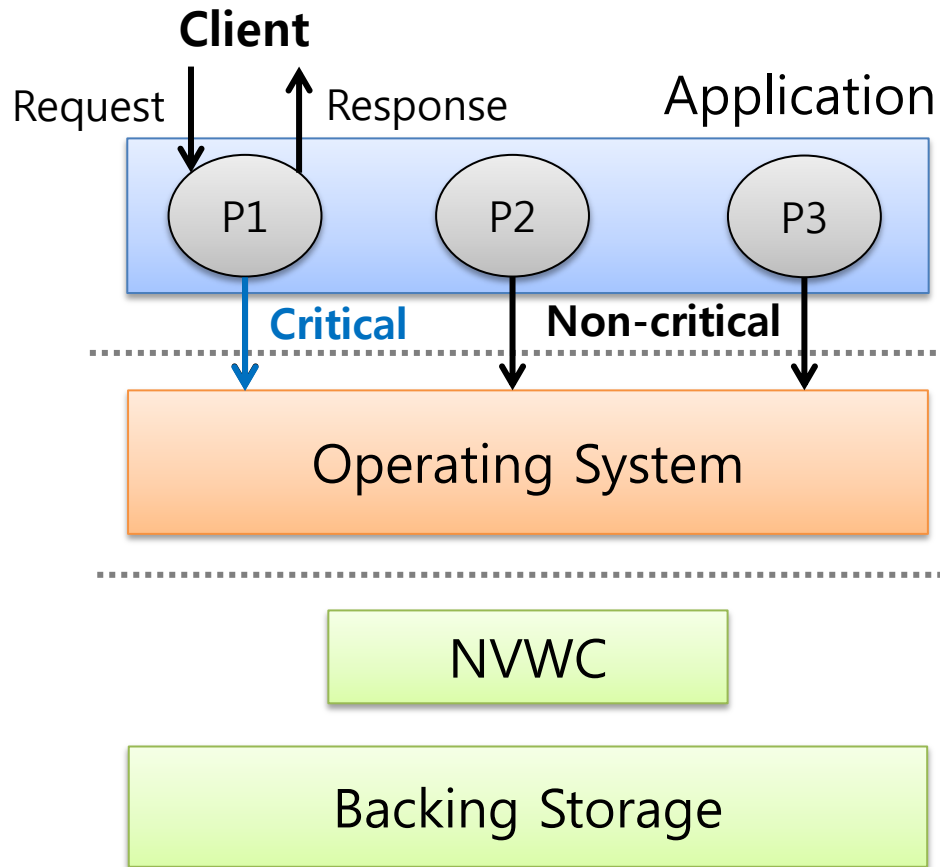
Criticality-Agnostic Contention

- Different write criticality



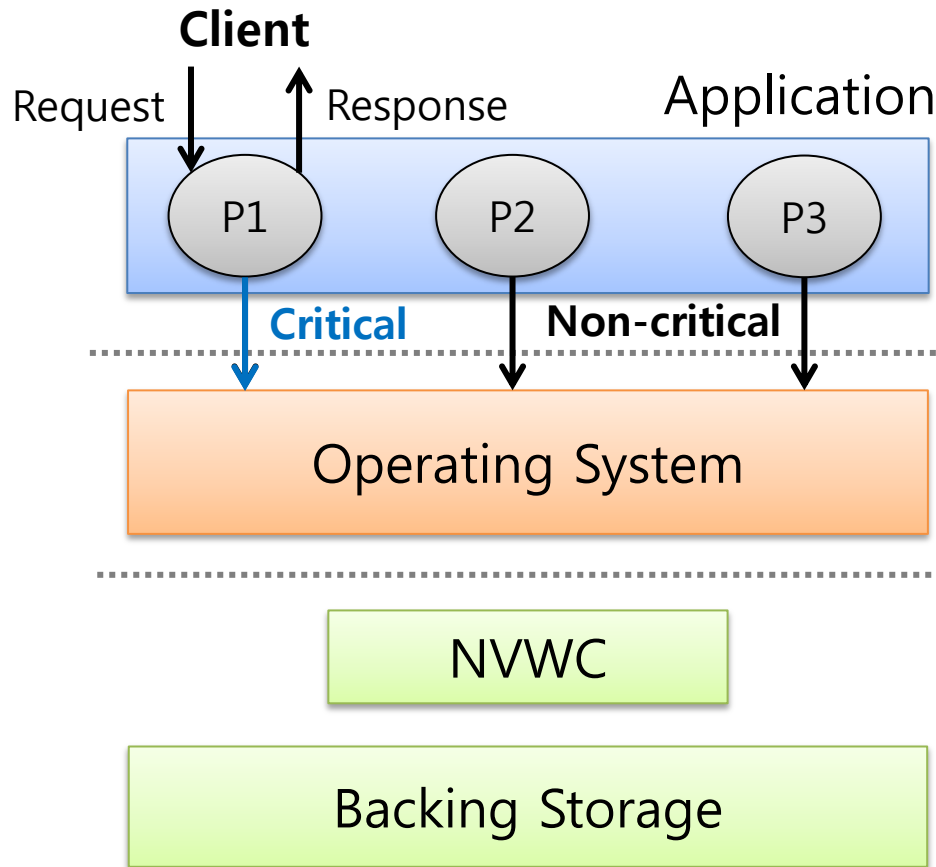
Criticality-Agnostic Contention

- Different write criticality



Criticality-Agnostic Contention

- Different write criticality

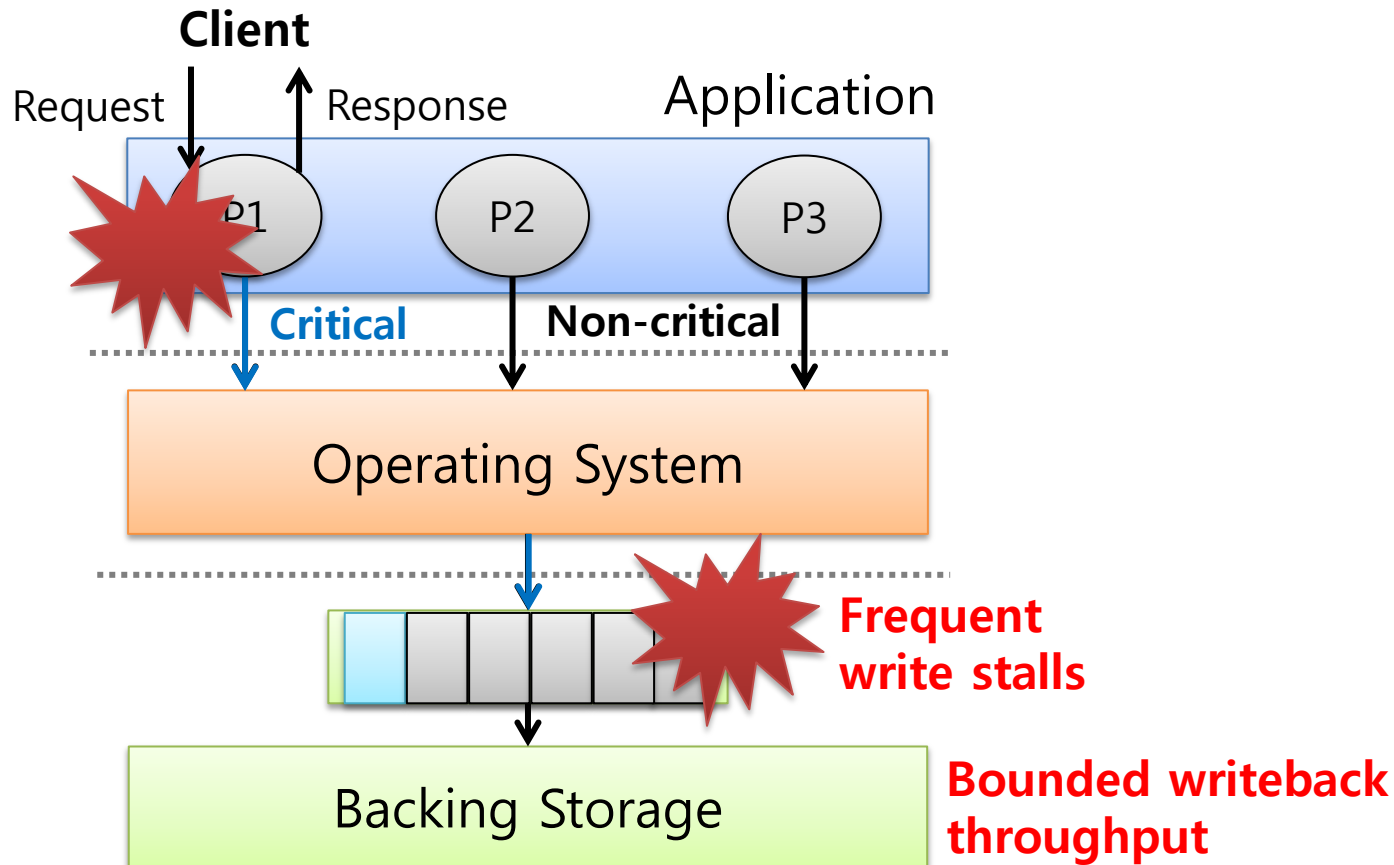


* Contentions

- Capacity contention
- Bandwidth contention

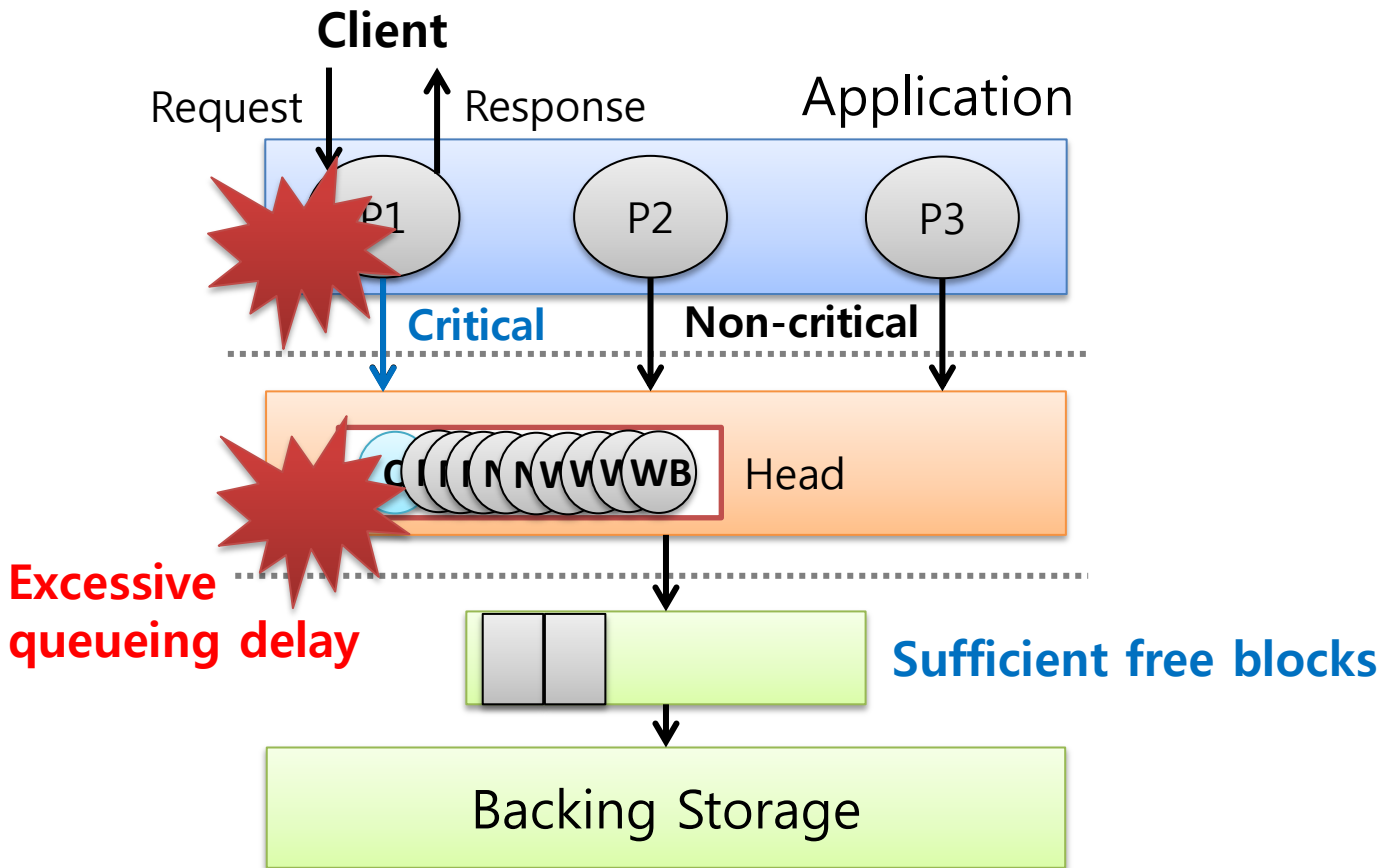
Criticality-Agnostic Contention

- Capacity contention



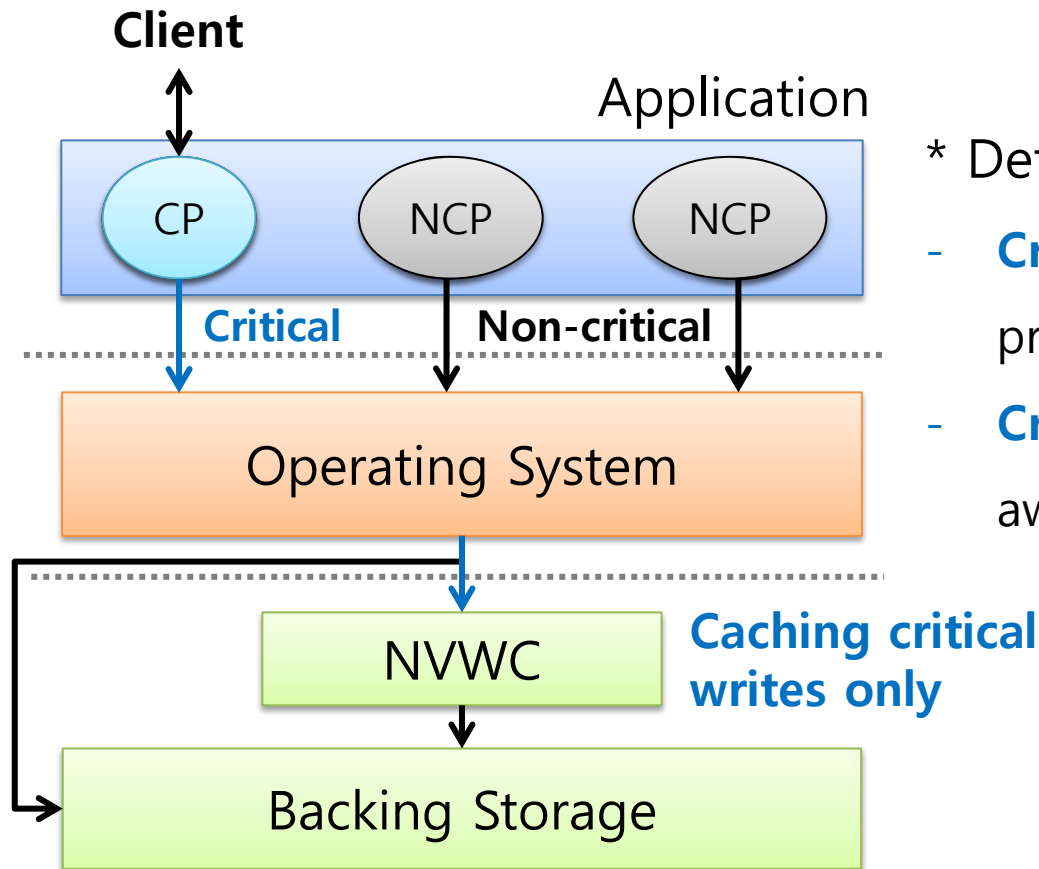
Criticality-Agnostic Contention

- Bandwidth contention



Our Approach

- Request-oriented caching policy



* Definitions

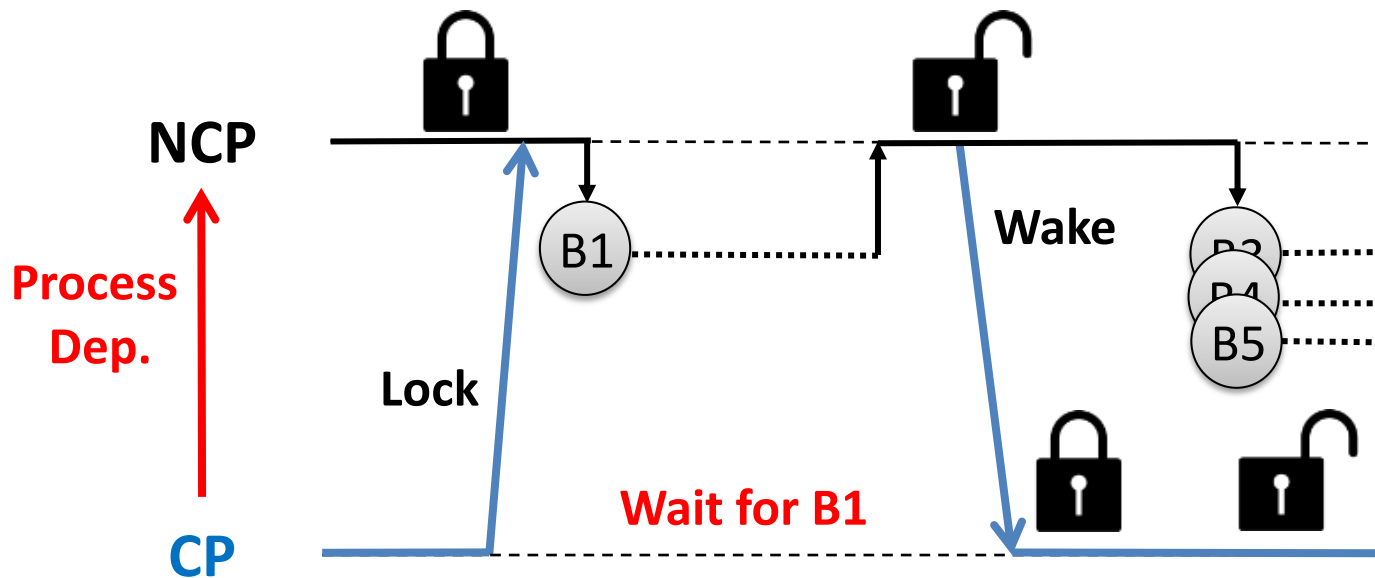
- **Critical process (CP)**: a process handling request
- **Critical write**: a write awaited by a critical proc.

Challenge

- How to accurately detect critical writes
- Types of critical write
 - Sync. writes from critical processes
 - **Dependency-induced** critical writes
 - Process dependency-induced
 - I/O dependency-induced


Dependency Problem

- Process dependency



Dependency Problem

- I/O dependency



* Example scenarios:

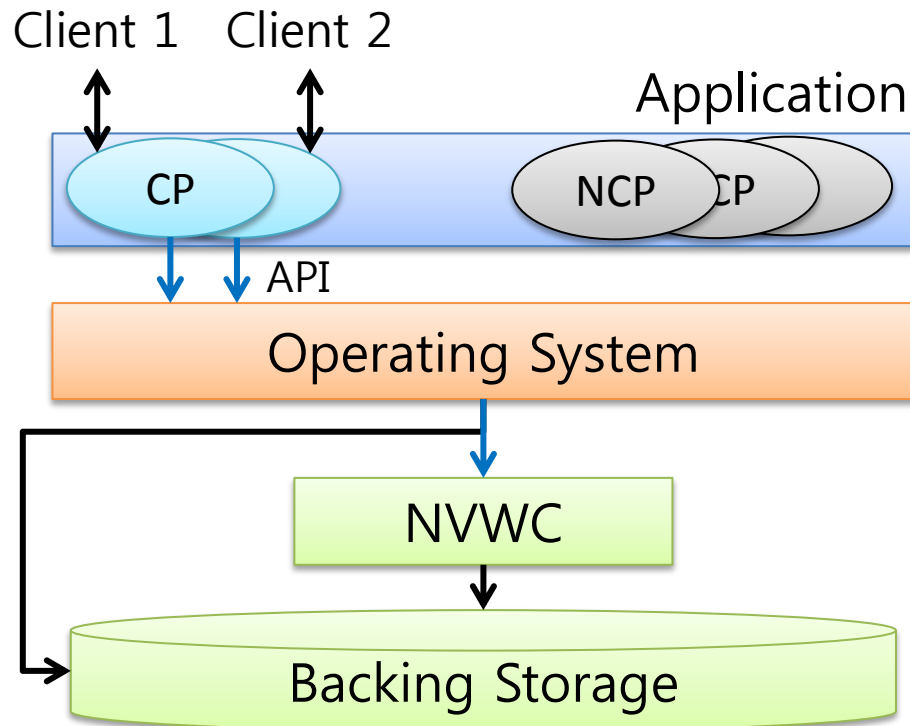
- CP **fsync()** to a block under writeback issued by NCP
- CP tries to **overwrite** fs journal buffer under writeback

Critical Write Detection

- **Critical process identification**
 - **Application-guided identification**

Critical Process Identification

- Application-guided identification



Critical Write Detection


- Critical process identification
 - Application-guided identification
- **Dependency resolution**
 - **Criticality inheritance protocols**
 - Process criticality inheritance
 - I/O criticality inheritance
 - Blocking object tracking

Critical Write Detection

- Critical process identification
 - Application-guided identification
- **Dependency resolution**
 - **Criticality inheritance protocols**
 - **Process criticality inheritance**
 - **I/O criticality inheritance**
 - **Blocking object tracking**


Criticality Inheritance Protocols

- Process criticality inheritance



Criticality Inheritance Protocols

- I/O criticality inheritance




Key issue:

caching the dependent write outstanding to disk w/o side effects

Evaluation

- Implementation on Linux 3.13 w/ FlashCache 3.1
- Application studies
 - PostgreSQL database

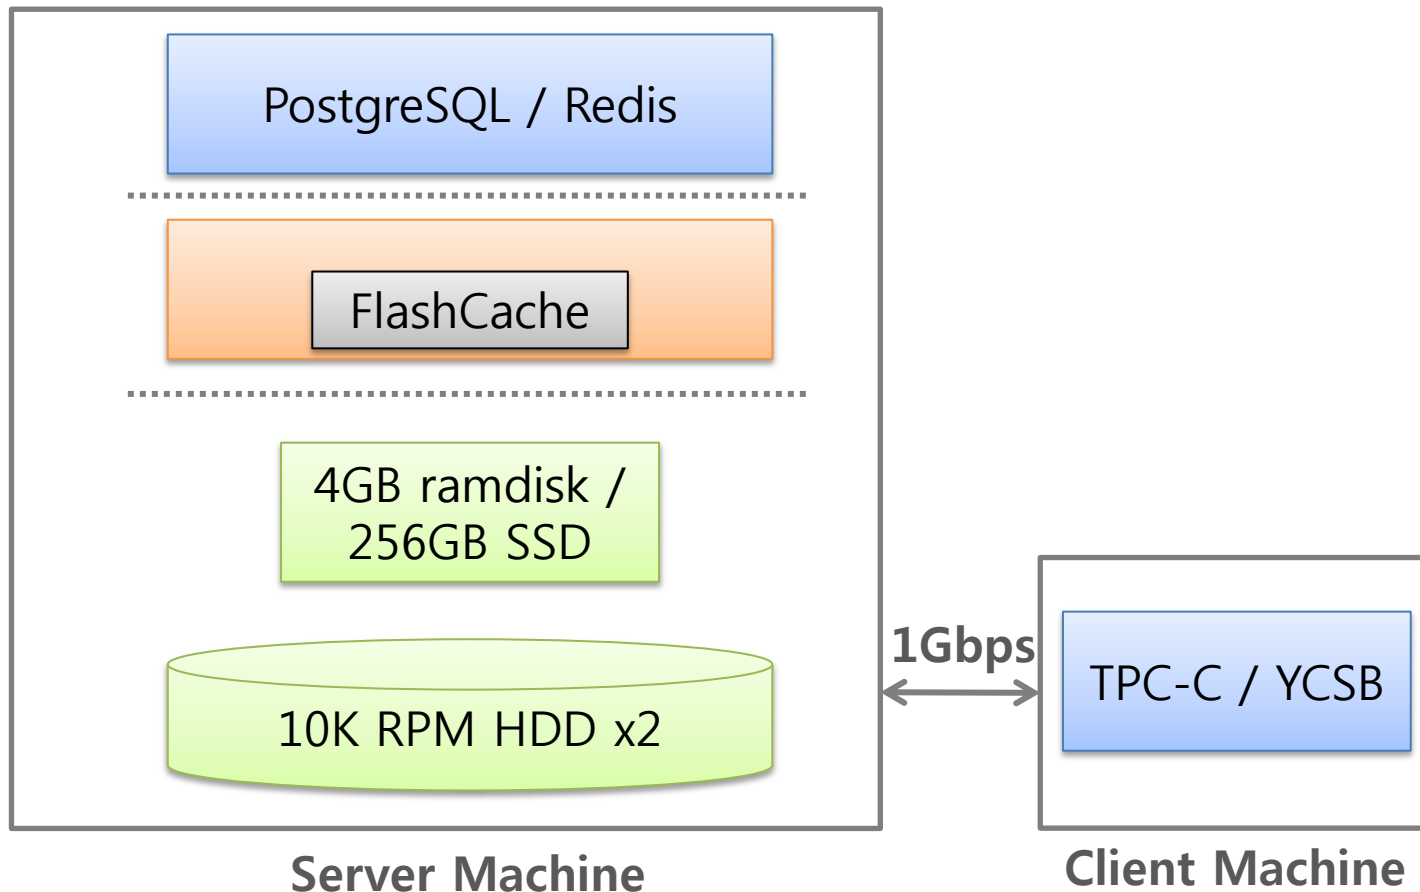


- Redis key-value store



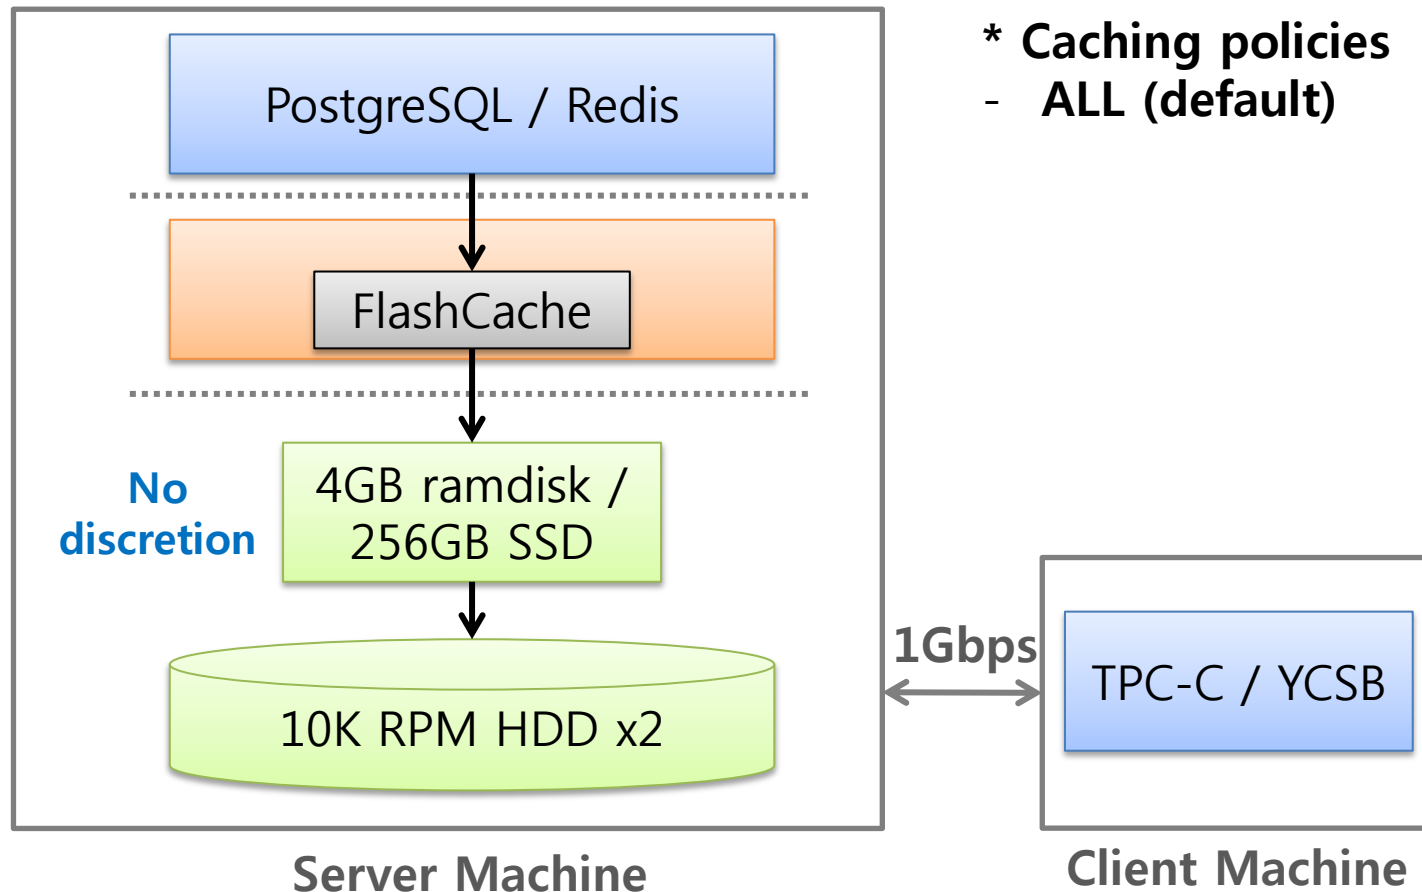
Evaluation

- Experimental setup



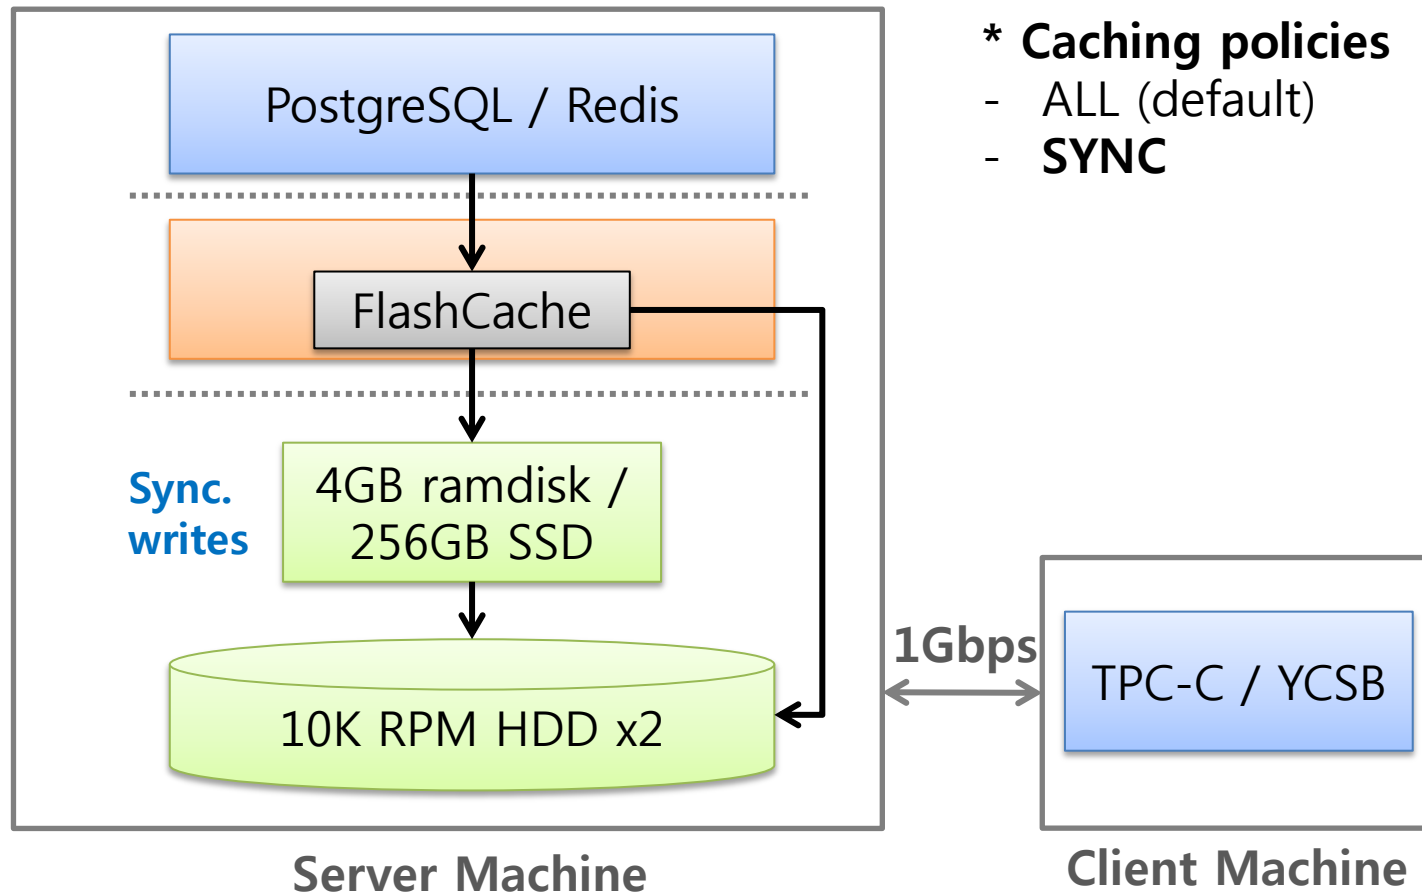
Evaluation

- Experimental setup



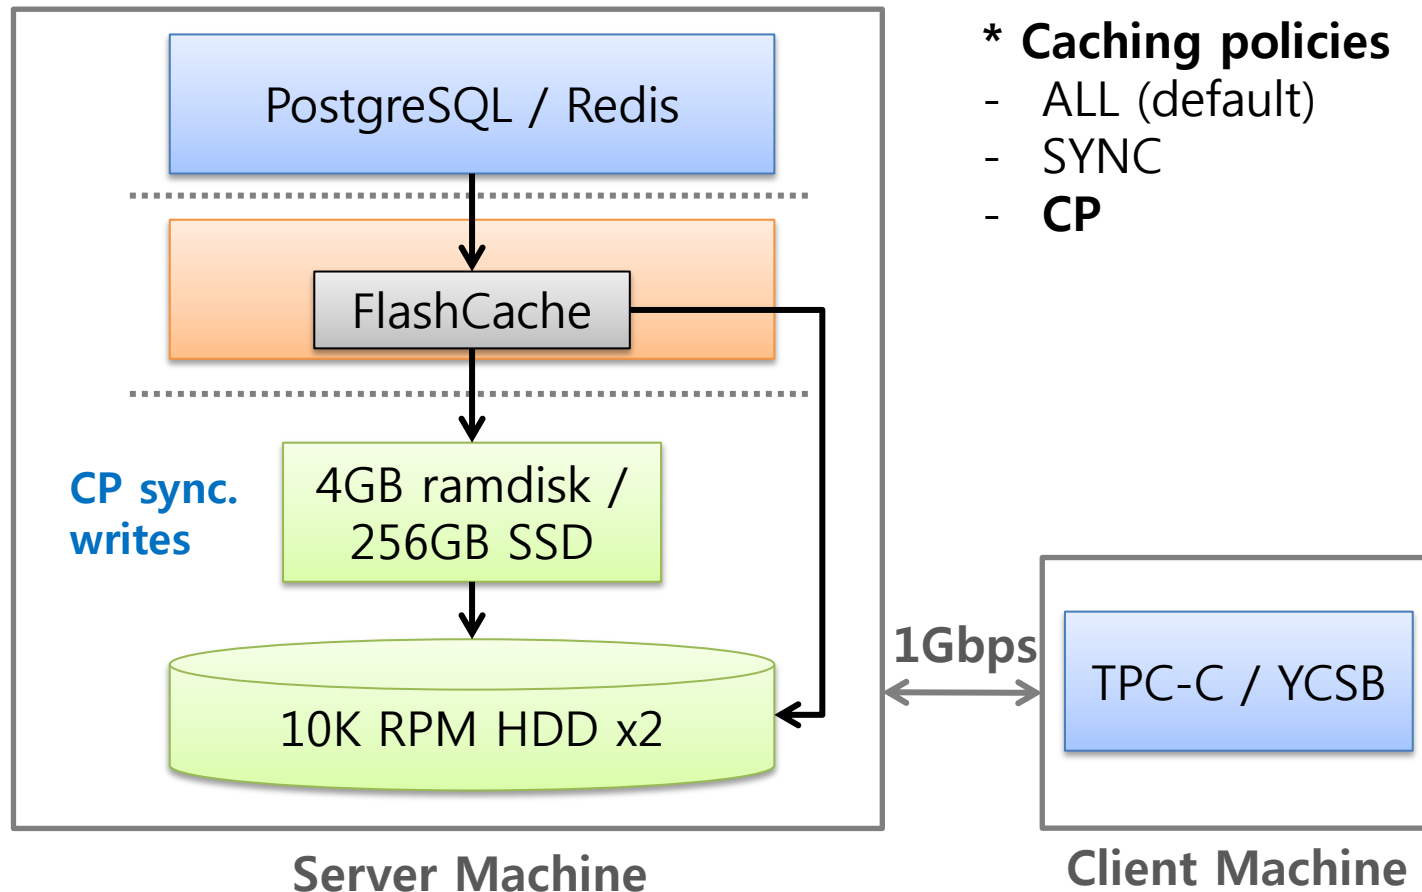
Evaluation

- Experimental setup



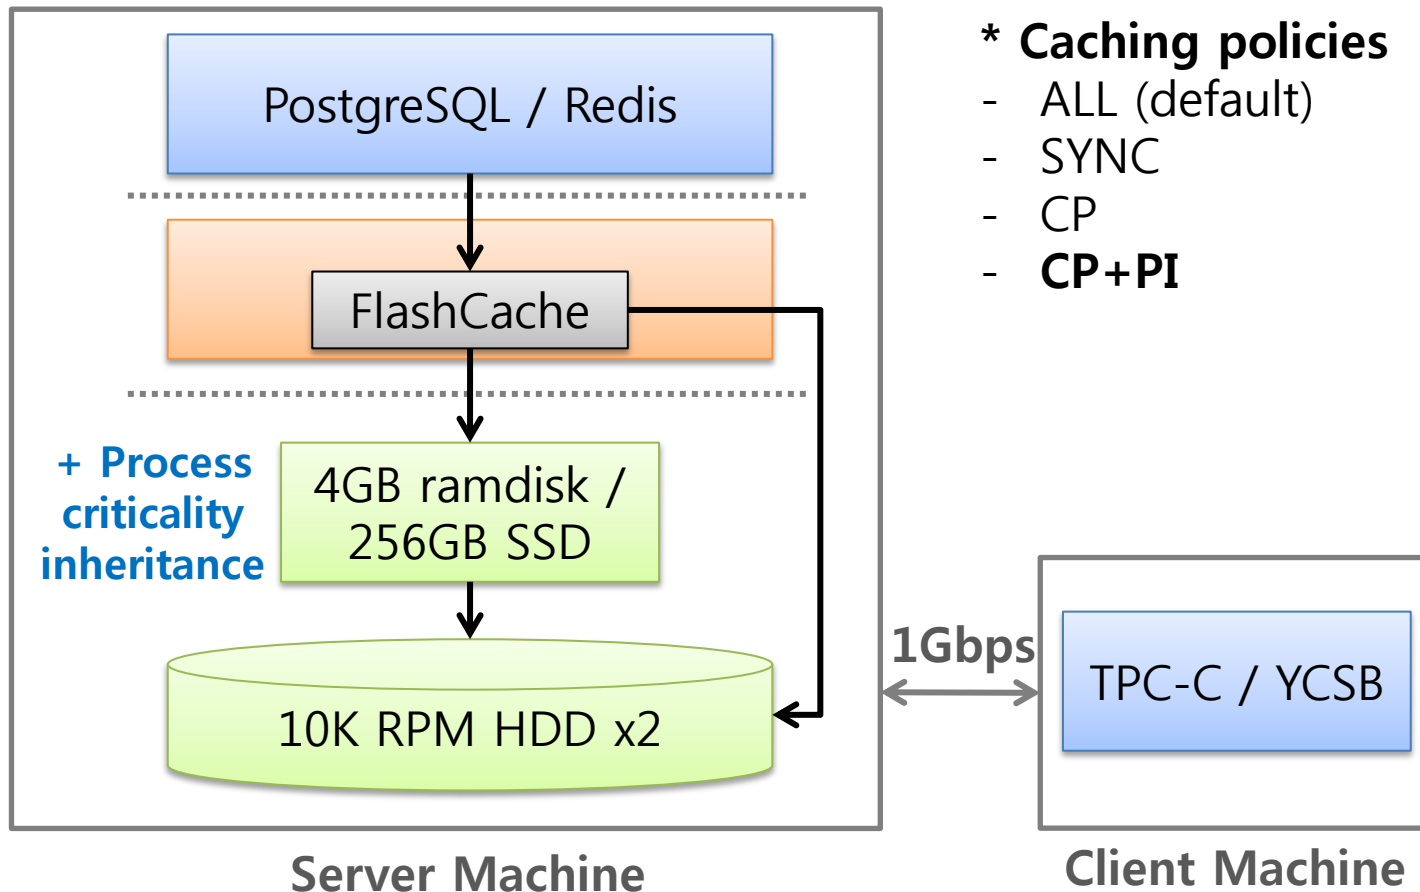
Evaluation

- Experimental setup



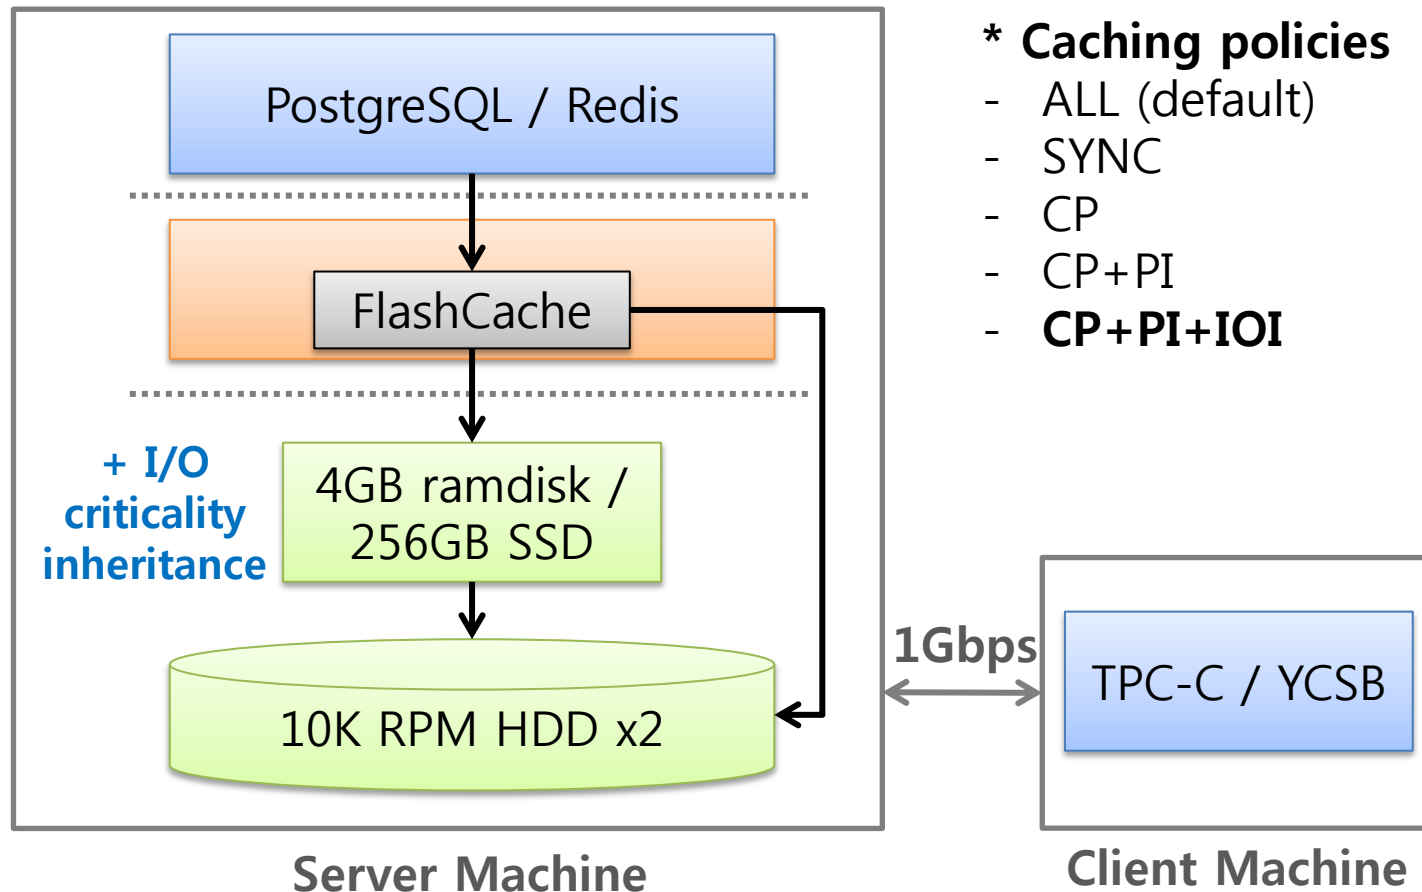
Evaluation

- Experimental setup



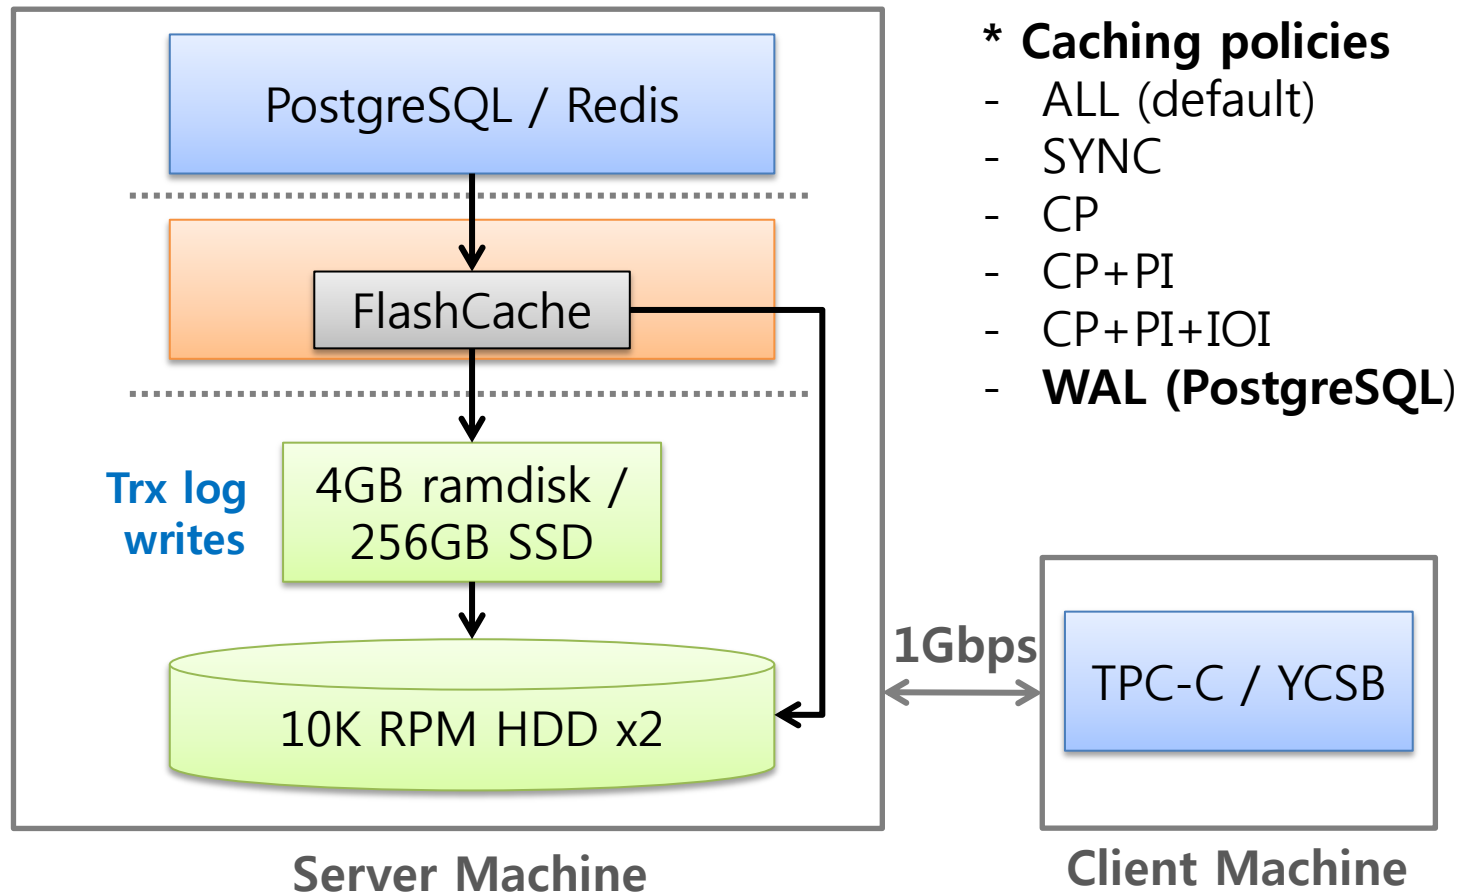
Evaluation

- Experimental setup



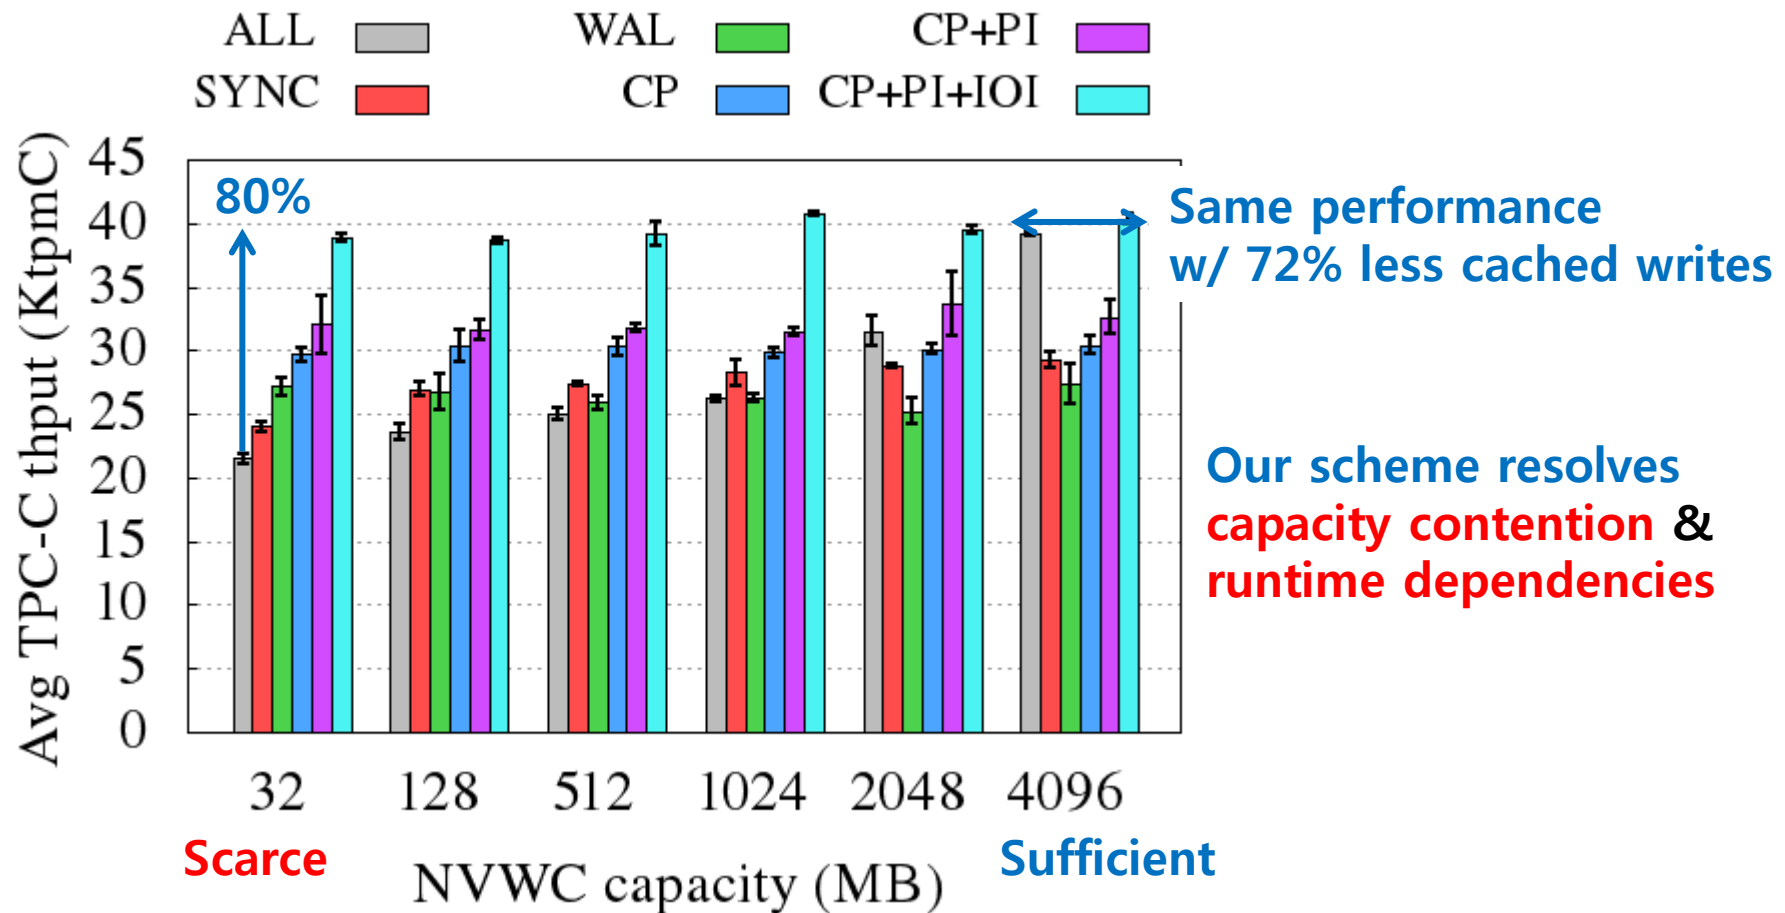
Evaluation

- Experimental setup



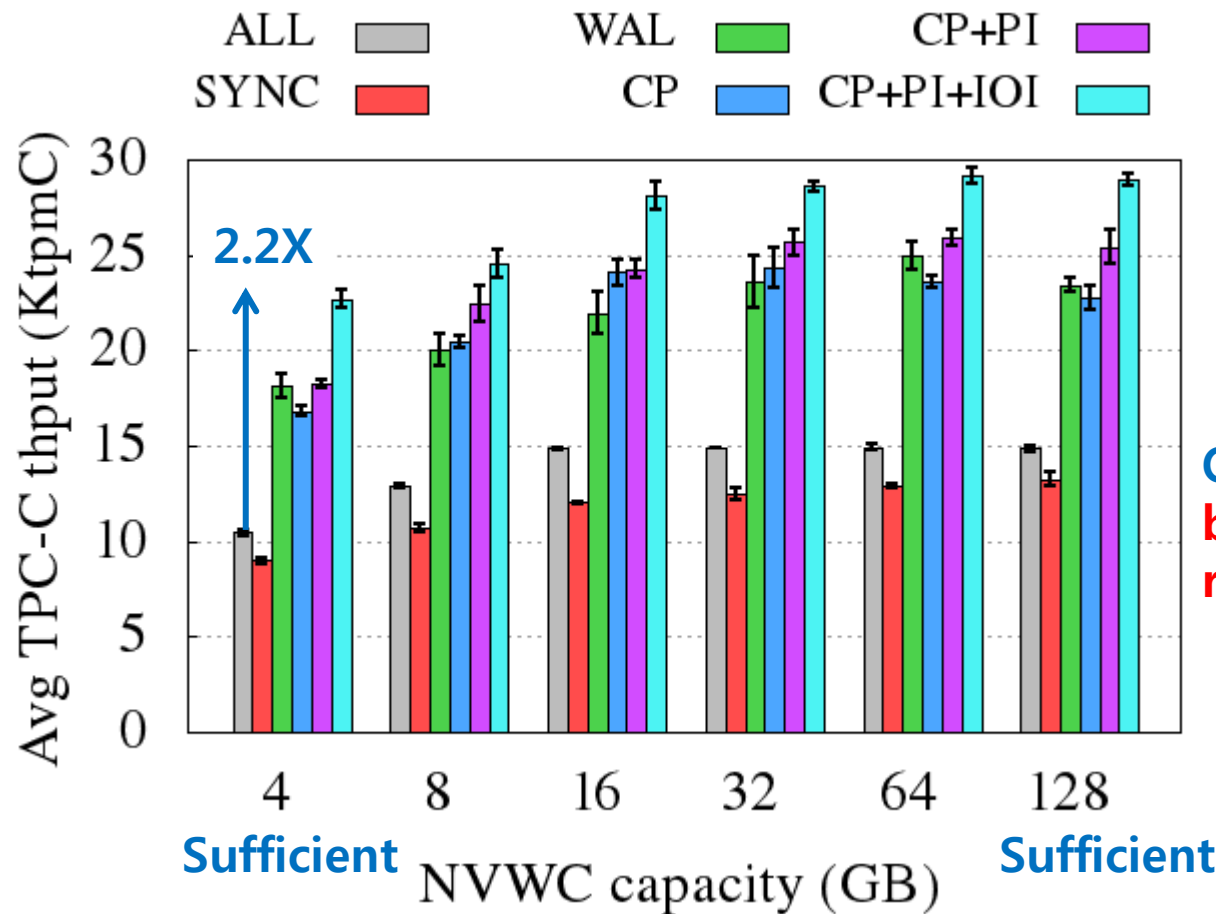
PostgreSQL Performance

- TPC-C workload w/ ramdisk



PostgreSQL Performance

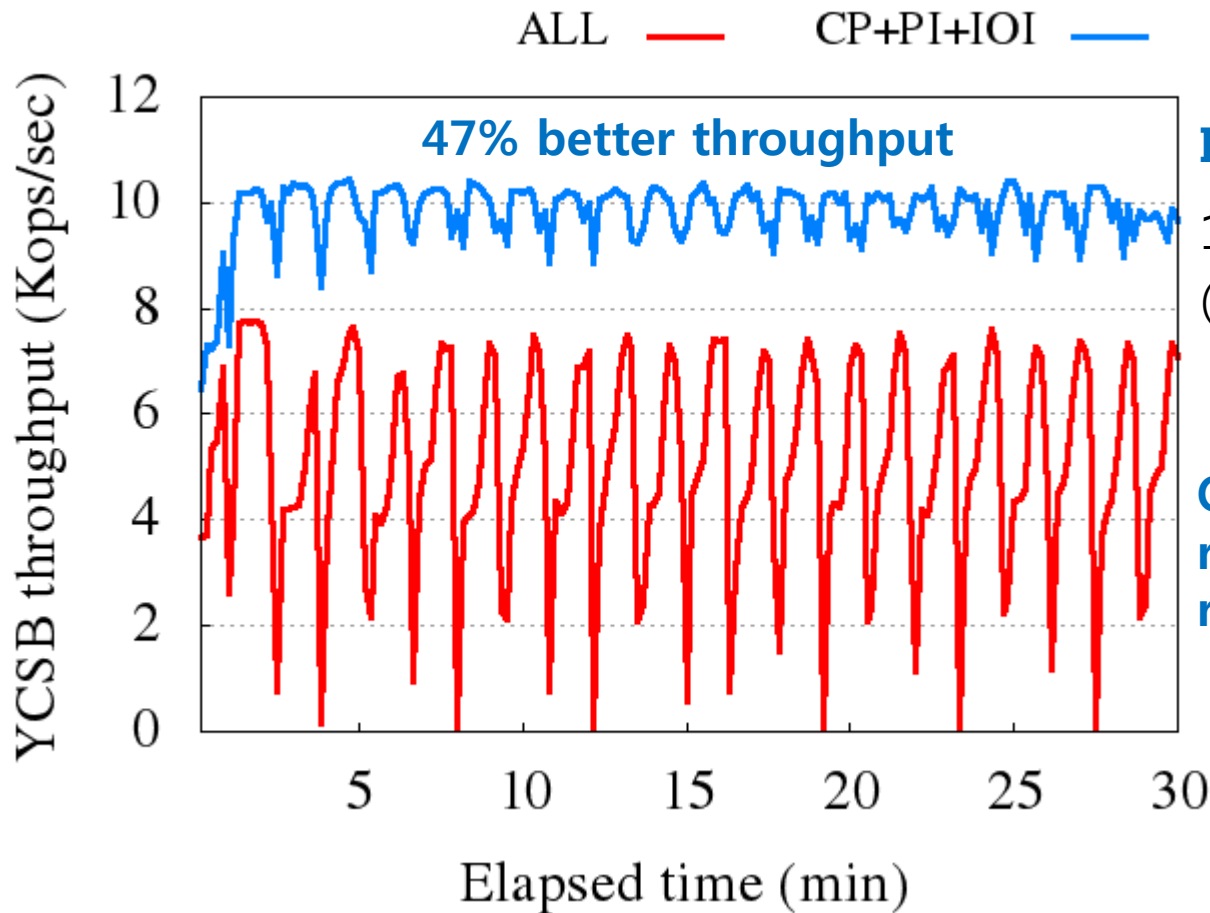
- TPC-C workload w/ SSD



Our scheme resolves
bandwidth contention &
runtime dependencies

Redis Performance

- Update-heavy workload w/ 16GB SSD



Improved tail latency

13X better @ 99.9th %ile
(50ms vs. 649ms)

**Our scheme improves
request throughput &
request latency**

Conclusions

- Key observation
 - Each write has different performance-criticality
- Request-oriented caching policy
 - Solely utilizes NVWC for application performance
 - **Improves** performance while **reducing** cached writes
- Future work
 - System-level critical process identification
 - Application to user-interactive environments

Thank You!

- Questions and comments
- Contact
 - sw.kim@skku.edu