

# 利用MySQL 5.7新特性做SQL优化

吴炳锡 - 知数堂培训联合创始人  
2017.3

# 关于我

- 知数堂培训联合创始人， MySQL中国用户组发起人之一及主席团成员
- QQ/微信： 82565387
- 10年+MySQL从业经验， 专注MySQL架构， 多IDC设计， 高可用， 高性能相关技术分析及实践， 目前从事MySQL培训及企业MySQL相关服务。

# agenda

- 怎么开展SQL优化
- 传统模型下做SQL优化需要做的工作
- 在MySQL 5.7下如何开展SQL优化
  - 利用Sys库发现需要优化的SQL
  - 利用Query Rewrite Plugin进行优化测试
  - 利用HINT改写优化及SQL执行超时保护
- 总结

- 事实上，MySQL 5.7已经不新了，2016年已发布8.0版本
- 2013.4.23发布MySQL 5.7.1，已经将近4年了
- 每年至少3个版本，2016年突然发飙，发布了7个版本
- 官方号称比5.6快3倍以上
- MySQL 5.7 让MySQL DBA有了更多的调优机会
- 截止目前，最新版本5.7.17，仍在持续完善中.....

# 如何开展SQL优化

占用时间最长的SQL

执长时间最多的SQL

业务高峰期的SQL





# 如何开展SQL优化

- **SQL优化的目标：**
  - 让系统跑的平稳
  - 有更高的业务处理能力
  - 更短的业务响应能力



# 传统模型下做SQL优化需要做的工作



## 存在问题

1. 开发需要提供SQL List列表
2. 对DBA对SQL有没有Where条件，索引使用合理不合理审查
3. 慢日志收集分析，一个月后，看惯了那个老脸也没新鲜了

# 传统模型下做SQL优化需要做的工作

## 获取SQL List优化

- 使用general log
- 把long\_query\_time=0
- 利用tcpdump类工具进行抓包
- 等等方法

**需要一个前提：**

能登录上这个服务器才行！

RDS，云上的DB怎么办？



# MySQL 5.7下开展 SQL优化

一个真正的DB，有DBA出手的机会




## MySQL 5.7: Connections per Second

1.7x Faster than MySQL 5.6

2.5x Faster than MySQL 5.5

67,000 Connections/Sec



Intel(R) Xeon(R) CPU E7-4860 v2 @ 2.3 GHz  
4 sockets x 10 cores-HT (80 CPU threads)  
2.3 GHz, 512 GB RAM  
Oracle Linux 6.5



Copyright © 2013, Oracle and/or its affiliates. All rights reserved.

14

1

# MySQL 5.7下开展SQL优化

- 针对SQL优化上改善
  - 更严谨的SQL约束
  - explain for connetion 1024
  - json数据类型
  - 多源复制，支持多个Triger
  - 优化器HINT提示
  - ServerSide执行超时

MySQL 5.7  
更完善  
更高效SQL支持  
等待着你来发现

# Sys Schema

- 从SYS Schema可以快速获取很多
  - 锁等待
  - 内存分配
  - SQL统计

| statement         | total    |
|-------------------|----------|
| Execute           | 16334741 |
| commit            | 642667   |
| rollback          | 2077     |
| Prepare           | 1050     |
| select            | 18       |
| Close stmt        | 350      |
| stmt              | 12       |
| set               | 4        |
| call_procedure    | 5        |
| show_variables    | 1        |
| jump_if_not       | 4266     |
| show_tables       | 5        |
| freturn           | 1040     |
| set_option        | 30       |
| show_table_status | 1        |
| Init DB           | 3        |
| Quit              | 11       |
| show_processlist  | 1        |

18 rows in set (0.00 sec)

参考：MySQL 5.7系列之sys schema

sys.user\_summary\_by\_statement\_type

# 利用performance\_schema

```
root@localhost:mysql3306.sock [sys]>select * from statement_analysis limit 10\G;
```

- sys.statement\_analysis
- CALL  
sys.ps\_setup\_enable\_instrument('statement');

```
      query: SELECT `c` FROM `sbtest1` WHERE `id` = ?
      db: bench
      full_scan:
      exec_count: 23785
      err_count: 0
      warn_count: 0
      total_latency: 8.15 s
      max_latency: 43.95 ms
      avg_latency: 342.57 us
      lock_latency: 602.96 ms
      rows_sent: 23785
      rows_sent_avg: 1
      rows_examined: 23785
      rows_examined_avg: 1
      rows_affected: 0
      rows_affected_avg: 0
      tmp_tables: 0
      tmp_disk_tables: 0
      rows_sorted: 0
      sort_merge_passes: 0
      digest: 80295d1d2720d4515b05d648e8caa82f
      first_seen: 2017-02-21 17:21:54
      last_seen: 2017-02-21 17:25:53
*****
```



# 利用performance\_schema获取

```
SELECT CONCAT_WS( '# Time: ', date_format(CURDATE(), '%y%m%d'),  
' TIME_FORMAT(NOW(6), '%H:%i:%s.%f'), '\n', '# User@Host:  
' , t.PROCESSLIST_USER, '[' , t.PROCESSLIST_USER, ' ] @ ', PROCESSLIST_HOST, ' [] Id:  
' , t.PROCESSLIST_ID, '\n', '# Schema: ', CURRENT_SCHEMA, ' Last_errno: ', MYSQL_ERRNO, '  
' , '\n', '# Query_time: ', ROUND(s.TIMER_WAIT / 1000000000000, 6), ' Lock_time:  
' , ROUND(s.LOCK_TIME / 1000000000000, 6), ' Rows_sent: ', ROWS_SENT, '  
Rows_examined: ', ROWS_EXAMINED, ' Rows_affected: ', ROWS_AFFECTED, '\n', '#  
Tmp_tables: ', CREATED_TMP_TABLES, ' Tmp_disk_tables: ', CREATED_TMP_DISK_TABLES, '  
' , '\n', '# Full_scan: ', IF(SELECT_SCAN=0, 'No', 'Yes'), ' Full_join:  
' , IF(SELECT_FULL_JOIN=0, 'No', 'Yes'), ' Tmp_table:  
' , IF(CREATED_TMP_TABLES=0, 'No', 'Yes'), ' Tmp_table_on_disk:  
' , IF(CREATED_TMP_DISK_TABLES=0, 'No', 'Yes'), '\n' , t.PROCESSLIST_INFO, ';' ) FROM  
performance_schema.events_statements_history s JOIN performance_schema.threads t  
using(thread_id) WHERE t.TYPE = 'FOREGROUND' AND t.PROCESSLIST_INFO IS NOT  
NULL AND t.PROCESSLIST_ID != connection_id() ORDER BY t.PROCESSLIST_TIME desc;
```



# 利用performance\_schema获取

```
# User@Host: wubx[wubx] @ 192.168.11.100 [] Id: 24
# Schema: bench Last_errno: 0
# Query_time: 0.000130 Lock_time: 0.000021 Rows_sent: 1 Rows_examined: 1 Rows_affected: 0
# Tmp_tables: 0 Tmp_disk_tables: 0
# Full_scan: No Full_join: No Tmp_table: No Tmp_table_on_disk: No
SELECT c FROM sbtest5 WHERE id=53561;
|
| # Time: 170221 17:44:50.414066
# User@Host: wubx[wubx] @ 192.168.11.100 [] Id: 24
# Schema: bench Last_errno: 0
# Query_time: 0.000114 Lock_time: 0.000019 Rows_sent: 1 Rows_examined: 1 Rows_affected: 0
# Tmp_tables: 0 Tmp_disk_tables: 0
# Full_scan: No Full_join: No Tmp_table: No Tmp_table_on_disk: No
SELECT c FROM sbtest5 WHERE id=53561;
|
+-----+
|
+-----+
40 rows in set (0.02 sec)
```

# MySQL 5.7下开展SQL优化

- 数据就在那里，也没额外开销 MySQL 5.6.6 后默认打开
- 生产环境中也可以放心使用
- 对于Sys库的一个统计计算就是Triger有视图的使用，开销很小，如果不需要可以禁掉。



# Query Rewrite Plugin

从容面对



# Query Rewrite Plugin介绍

- 在MySQL 5.7.6以后官方引入一个Query Rewrite plugin，可以在Server端收到SQL后，进行改写并执行。
  - 只针对标准的SELECT语句工作，不能对视图定义及存储过程中SELECT语句改写
  - 改写规则记录在内存中，实际对应到：query\_rewrite库下的rewrite\_rules这个表
  - 利用query\_rewrite下的存储过程：load\_rewrite\_rules() / flush\_rewrite\_rules() 及DML语句来加载更改规则

# Query Rewrite Plugin使用

- 启用Query Rewrite Plugin

```
mysql -u root -p < $basedir/share/install_rewriter.sql  
SHOW GLOBAL VARIABLES LIKE 'rewriter_enabled';
```

| Variable_name    | Value |
|------------------|-------|
| rewriter_enabled | ON    |

- 在配置文件中添加

```
[mysqld]  
rewriter_enabled=ON
```



# Query Rewrite Plugin使用

- 利用子查询新特性优化Join

```
root@localhost [employees]>root@localhost [employees]>SELECT count(distinct emp_no) FROM employees.employees INNER JOIN employees.salaries ON employees.emp_no=employees.salaries.emp_no WHERE DATEDIFF(to_date, from_date) < 2000;
```

| count(distinct emp_no) |
|------------------------|
| 300024                 |

```
1 row in set (3.09 sec)
```

- 优化后样子

```
root@localhost [employees]>SELECT count(emp_no) FROM employees.employees WHERE emp_no IN ( SELECT emp_no FROM employees.salaries WHERE DATEDIFF(to_date, from_date)<2000);
```

| count(emp_no) |
|---------------|
| 300024        |

```
1 row in set (1.37 sec)
```

# Query Rewrite Plugin使用

- 利用Query Rewrite固化

```
INSERT INTO query_rewrite.rewrite_rules
(
pattern,
Replacement
)
VALUES
(
'SELECT count(distinct emp_no) FROM employees.employees INNER JOIN
employees.salaries USING(emp_no) WHERE DATEDIFF(to_date, from_date) < ?',
'SELECT count(emp_no) FROM employees.employees WHERE emp_no IN ( SELECT emp_no
FROM employees.salaries WHERE DATEDIFF(to_date, from_date) < ?) '
);
CALL query_rewrite.flush_rewrite_rules();
```

# ServerSide SQL执行超时

避免对DB过载

# max\_execution\_time

- MySQL 5.7.8后引入的参数
- 默认为0，不限制。如果在全局设置了这个参数，相当于引入全局的ServerSide超时
- 同时可以作用于Session中的`/*+ max_execution_time(N) */` #N单位毫秒
- **只对Select语句起作用，但对存储过程中的Select语句不起作用**

# max\_execution\_time

例子：

```
root@localhost [employees]>root@localhost [employees]>SELECT count(distinct emp_no) FROM employees.employees INNER JOIN employees.sa
DATEDIFF(to_date, from_date) < 2000;
+-----+
| count(distinct emp_no) |
+-----+
|           300024 |
+-----+
1 row in set (3.09 sec)
```

业务高峰期，某些SQL处理上，我希望给降级或是引入过载保护，让DB能撑过业务高峰，对于这个问题我们怎么处理？

pt-kill?



# max\_execution\_time

- 结合Query Rewrite Plugin 和HINT改写

```
INSERT INTO query_rewrite.rewrite_rules
(
pattern,
Replacement
)
VALUES
(
'SELECT count(distinct emp_no) FROM employees.employees INNER JOIN
employees.salaries USING(emp_no) WHERE DATEDIFF(to_date, from_date) < ?',
'SELECT /*+ MAX_EXECUTION_TIME(1000)*/ count(distinct emp_no) FROM
employees.employees INNER JOIN employees.salaries USING(emp_no) WHERE
DATEDIFF(to_date, from_date) < ?',
);

CALL query_rewrite.flush_rewrite_rules();
```

# max\_execution\_time

执行效果

```
root@localhost [employees]>SELECT count(distinct emp_no) FROM employees.employees INNER JOIN employees.salaries USING(emp_no) WHERE  
DATEDIFF(to_date, from_date) < 2000;  
ERROR 3024 (HY000): Query execution was interrupted, maximum statement execution time exceeded  
root@localhost [employees]>SELECT count(distinct emp_no) FROM employees.employees INNER JOIN employees.salaries USING(emp_no) WHERE  
DATEDIFF(to_date, from_date) < 2000;  
ERROR 3024 (HY000): Query execution was interrupted, maximum statement execution time exceeded
```

# HINT增强

SELECT /\*+ ... \*/ ...

INSERT /\*+ ... \*/ ...

REPLACE /\*+ ... \*/ ...

UPDATE /\*+ ... \*/ ...

DELETE /\*+ ... \*/ ...

SELECT /\*+ NO\_RANGE\_OPTIMIZATION(t3 PRIMARY, f2\_idx) \*/ \* FROM t3...;

SELECT /\*+ BKA(t1) NO\_BKA(t2) \*/ \* FROM t1 INNER JOIN t2 WHERE ...;

SELECT /\*+ NO\_ICP(t1, t2) \*/ \* FROM t1 INNER JOIN t2 WHERE ...;

SELECT /\*+ SEMIJOIN(FIRSTMATCH, LOOSESCAN) \*/ \* FROM t1 ...;

SELECT /\*+ NO\_ICP(t1) \*/ \* FROM t1 WHERE ...;



# 案例 排序过多

```
SELECT sl.sku,sl.positionCode,ps.positionLev,
SUM(CASE WHEN sk.skipQty > 0 THEN 0 ELSE CASE WHEN ps.positionLev = 3
      THEN (IFNULL(sl.quantity,0)-IFNULL(ll.lockedQty,0)+IFNULL(fl.frozenQty,0)) ELSE 0 END END) jhqty,
SUM(CASE WHEN sk.skipQty > 0 THEN 0 ELSE CASE WHEN ps.positionLev = 4
      THEN (IFNULL(sl.quantity,0)-IFNULL(ll.lockedQty,0)+IFNULL(fl.frozenQty,0)) ELSE 0 END END) bhqty,
SUM(CASE WHEN sk.skipQty > 0 THEN 0 ELSE CASE WHEN ps.positionLev = 5
      THEN (IFNULL(sl.quantity,0)-IFNULL(ll.lockedQty,0)+IFNULL(fl.frozenQty,0)) ELSE 0 END END) byqty,
SUM(CASE WHEN sk.skipQty > 0 THEN 0 ELSE CASE WHEN ps.positionLev = 6
      THEN (IFNULL(sl.quantity,0)-IFNULL(ll.lockedQty,0)+IFNULL(fl.frozenQty,0)) ELSE 0 END END) ryqty
FROM wms_si_stock sl
JOIN wms_bd_product_information p ON p.sku=sl.sku
JOIN wms_bd_positions ps ON ps.warehouseIdentity=sl.warehouseIdentity AND ps.positionCode=sl.positionCode
LEFT JOIN (SELECT l.warehouseIdentity,l.sku,l.positionCode,l.isbox,l.boxNo,IFNULL(SUM(quantity),0) lockedQty
FROM wms_si_locked l GROUP BY l.warehouseIdentity,l.sku,l.positionCode,l.isbox,l.boxNo) ll
ON sl.warehouseIdentity=ll.warehouseIdentity AND sl.sku=ll.sku AND sl.positionCode=ll.positionCode AND
sl.isbox=ll.isbox AND sl.boxNo=ll.boxNo
LEFT JOIN (SELECT warehouseIdentity,sku,positionCode ,SUM(quantity) frozenQty FROM wms_so_frozen WHERE STATUS=1
AND 0>quantity AND checkPlanId !='' AND checkPlanId IS NOT NULL GROUP BY warehouseIdentity,sku,positionCode) fl
ON sl.warehouseIdentity=fl.warehouseIdentity AND sl.sku=fl.sku AND sl.positionCode=fl.positionCode
LEFT JOIN (SELECT warehouseIdentity,sku,positionCode,boxNo,SUM(quantity) skipQty FROM wms_si_skipped WHERE STATUS=0
and warehouseIdentity = '1'
GROUP BY warehouseIdentity,sku,positionCode,boxNo) sk
ON sk.warehouseIdentity=sl.warehouseIdentity
AND sk.sku=sl.sku
AND sk.positionCode=sl.positionCode AND sk.boxNo=sl.boxNo
where 1=1
and sl.warehouseIdentity = '1'
AND ps.positionLev != 7
GROUP BY sl.sku,sl.positionCode,ps.positionLev,sl.warehouseIdentity
\G;
```

引自： 知数堂 郑松华



# 案例 排序过多

```
***** 1. row *****
  id: 1
  select_type: PRIMARY
  table: ps
  type: ALL
possible_keys: idx_pos_pos,idx_pos_lev,idx_pos_posLev
  key: NULL
  key_len: NULL
  ref: NULL
  rows: 108416
  Extra: Using where; Using temporary; Using filesort
***** 2. row *****
  id: 1
  select_type: PRIMARY
  table: sl
  type: ref
possible_keys: idx_stock_all,idx_st_skuP,idx_st_posCd
  key: idx_st_posCd
  key_len: 152
  ref: wms.ps.positionCode
  rows: 1
  Extra: Using where
***** 3. row *****
  id: 1
  select_type: PRIMARY
  table: p
  type: ref
possible_keys: idx_proif_sku
  key: idx_proif_sku
  key_len: 153
  ref: wms.sl.sku
  rows: 1
  Extra: Using index
***** 4. row *****
  id: 1
  select_type: PRIMARY
  table: <derived2>
  type: ref
possible_keys: <auto_key0>
  key: <auto_key0>
  key_len: 368
  ref: wms.ps.warehouseIdentity,wms.sl.sku,wms.ps.positionCode,w
  rows: 10
  Extra: NULL
```

```
  key_len: 305
  ref: wms.ps.warehouseIdentity,wms.sl.sku,wms.ps.positionCode
  rows: 10
  Extra: NULL
***** 6. row *****
  id: 1
  select_type: PRIMARY
  table: <derived4>
  type: ref
possible_keys: <auto_key0>
  key: <auto_key0>
  key_len: 367
  ref: wms.ps.warehouseIdentity,wms.sl.sku,wms.ps.positionCode,wms.sl.boxNo
  rows: 2
  Extra: NULL
***** 7. row *****
  id: 4
  select_type: DERIVED
  table: wms_si_skipped
  type: ref
possible_keys: idx_skip_sku,idx_skip_stat
  key: idx_skip_stat
  key_len: 2
  ref: const
  rows: 22
  Extra: Using where; Using temporary; Using filesort
***** 8. row *****
  id: 3
  select_type: DERIVED
  table: wms_so_frozen
  type: ALL
possible_keys: idx_fz_skuP
  key: NULL
  key_len: NULL
  ref: NULL
  rows: 37532
  Extra: Using where; Using temporary; Using filesort
***** 9. row *****
  id: 2
  select_type: DERIVED
  table: l
  type: ALL
possible_keys: NULL
  key: NULL
  key_len: NULL
  ref: NULL
  rows: 15957
```



# 案例 排序过多

```
SELECT s.sku,s.positionCode,s.positionLev,
SUM(CASE WHEN s.skipQty > 0 THEN 0 ELSE CASE WHEN s.positionLev = 3
      THEN (IFNULL(s.quantity,0)-IFNULL(s.lockedQty,0)+IFNULL(s.frozenQty,0)) ELSE 0 END END) jhqty,
SUM(CASE WHEN s.skipQty > 0 THEN 0 ELSE CASE WHEN s.positionLev = 4
      THEN (IFNULL(s.quantity,0)-IFNULL(s.lockedQty,0)+IFNULL(s.frozenQty,0)) ELSE 0 END END) bhqty,
SUM(CASE WHEN s.skipQty > 0 THEN 0 ELSE CASE WHEN s.positionLev = 5
      THEN (IFNULL(s.quantity,0)-IFNULL(s.lockedQty,0)+IFNULL(s.frozenQty,0)) ELSE 0 END END) byqty,
SUM(CASE WHEN s.skipQty > 0 THEN 0 ELSE CASE WHEN s.positionLev = 6
      THEN (IFNULL(s.quantity,0)-IFNULL(s.lockedQty,0)+IFNULL(s.frozenQty,0)) ELSE 0 END END) ryqty
from (
SELECT sl.sku,sl.positionCode,ps.positionLev,sl.warehouseIdentity,sl.quantity
(SELECT IFNULL(SUM(quantity),0)
FROM wms_si_locked ll
ON sl.sku=ll.sku AND sl.positionCode=ll.positionCode and sl.warehouseIdentity=ll.warehouseIdentity AND
sl.isbox=ll.isbox AND sl.boxNo=ll.boxNo) lockedQty,
(SELECT SUM(fl.quantity) frozenQty
FROM wms_so_frozen fl WHERE fl.STATUS=1
AND 0>fl.quantity AND fl.checkPlanId !='' AND fl.checkPlanId IS NOT NULL
and sl.sku=fl.sku AND sl.positionCode=fl.positionCode AND sl.warehouseIdentity=fl.warehouseIdentity )frozenQty,
(SELECT SUM(sk.quantity) FROM wms_si_skipped sk
WHERE sk.STATUS=0
and sk.warehouseIdentity = '1' and sk.warehouseIdentity=sl.warehouseIdentity AND sk.sku=sl.sku AND
sk.positionCode=sl.positionCode AND sk.boxNo=sl.boxNo) skipQty
FROM wms_si_stock sl
JOIN wms_bd_product_information p ON p.sku=sl.sku
JOIN wms_bd_positions ps ON ps.warehouseIdentity=sl.warehouseIdentity AND ps.positionCode=sl.positionCode
where 1=1
and sl.warehouseIdentity = '1'
AND ps.positionLev != 7
) s
GROUP BY s.sku,s.positionCode,s.positionLev,s.warehouseIdentity
```

引自： 知数堂 郑松华

# 案例 排序过多

### 修改后的执行计划：

| id | select_type        | table      | type | possible_keys                           | key           | key_len | ref                              | rows   | Extra                           |
|----|--------------------|------------|------|---|---------------|---------|----------------------------------|--------|---------------------------------|
| 1  | PRIMARY            | <derived2> | ALL  | NULL                                    | NULL          | NULL    | NULL                             | 54224  | Using temporary; Using filesort |
| 2  | DERIVED            | ps         | ALL  | idx_pos_pos,idx_pos_lev,idx_pos_posLev  | NULL          | NULL    | NULL                             | 108416 | Using where                     |
| 2  | DERIVED            | sl         | ref  | idx_stock_all,idx_st_skuP,idx_st_posCd  | idx_st_posCd  | 152     | wms.ps.positionCode              | 1      | Using where                     |
| 2  | DERIVED            | p          | ref  | idx_proif_sku                           | idx_proif_sku | 153     | wms.sl.sku                       | 1      | Using index                     |
| 5  | DEPENDENT SUBQUERY | sk         | ref  | idx_skip_box,idx_skip_sku,idx_skip_stat | idx_skip_box  | 214     | wms.sl.boxNo,wms.sl.positionCode | 1      | Using where                     |
| 4  | DEPENDENT SUBQUERY | fl         | ref  | idx_fz_skuP,idx_fz_posCd                | idx_fz_skuP   | 304     | wms.sl.sku,wms.sl.positionCode   | 5      | Using where                     |
| 3  | DEPENDENT SUBQUERY | ll         | ref  | idx_lk_skuP,idx_lk_posCd,idx_lock_box   | idx_lk_skuP   | 304     | wms.sl.sku,wms.sl.positionCode   | 1      | Using where                     |

之前的SQL 运行了 9000多秒 修改之后运行了7s

# SQL优化三板斧

- 确认join字段类型一致和字符一致
- 去除不必要的排序及临时表
- 前置运算移到Where条件后增加过滤



# 总结

- 1.需要有全局观，知道自已的业务使用了MySQL的功能
- 2.需要要去了解一下MySQL有那些特性及功能能被所用
- 3.对于新特性
  - 如果是一个重大的更新，如果慎重
  - 对于Tips的功能，可以轻松使用



# 总结

- SQL优化的几个层面
  - 第一层面
    - 通过基础的书集学习SQL，正确编写SQL，通过索引优化（看书可以搞定，理解索引对优化的帮助）
  - 第二层面
    - SQL改写能力，有集合的意识，借助于优化器输出改写SQL，能搞明白Join的顺序，搞明白 Access 和filter的区别，通过改写得到正确的结果（通过练习就可达到）
  - 第三层面
    - 有一定的全局意识，并且针对看到的SQL，不使用优化器输出，也可以进行优化（已经对group by，join 及不同版本的优化器特性熟记于心），SQL优化已经成为工作中。（需要有系统的方法论）

谢谢， 希望有所帮助

